

University of
Kent

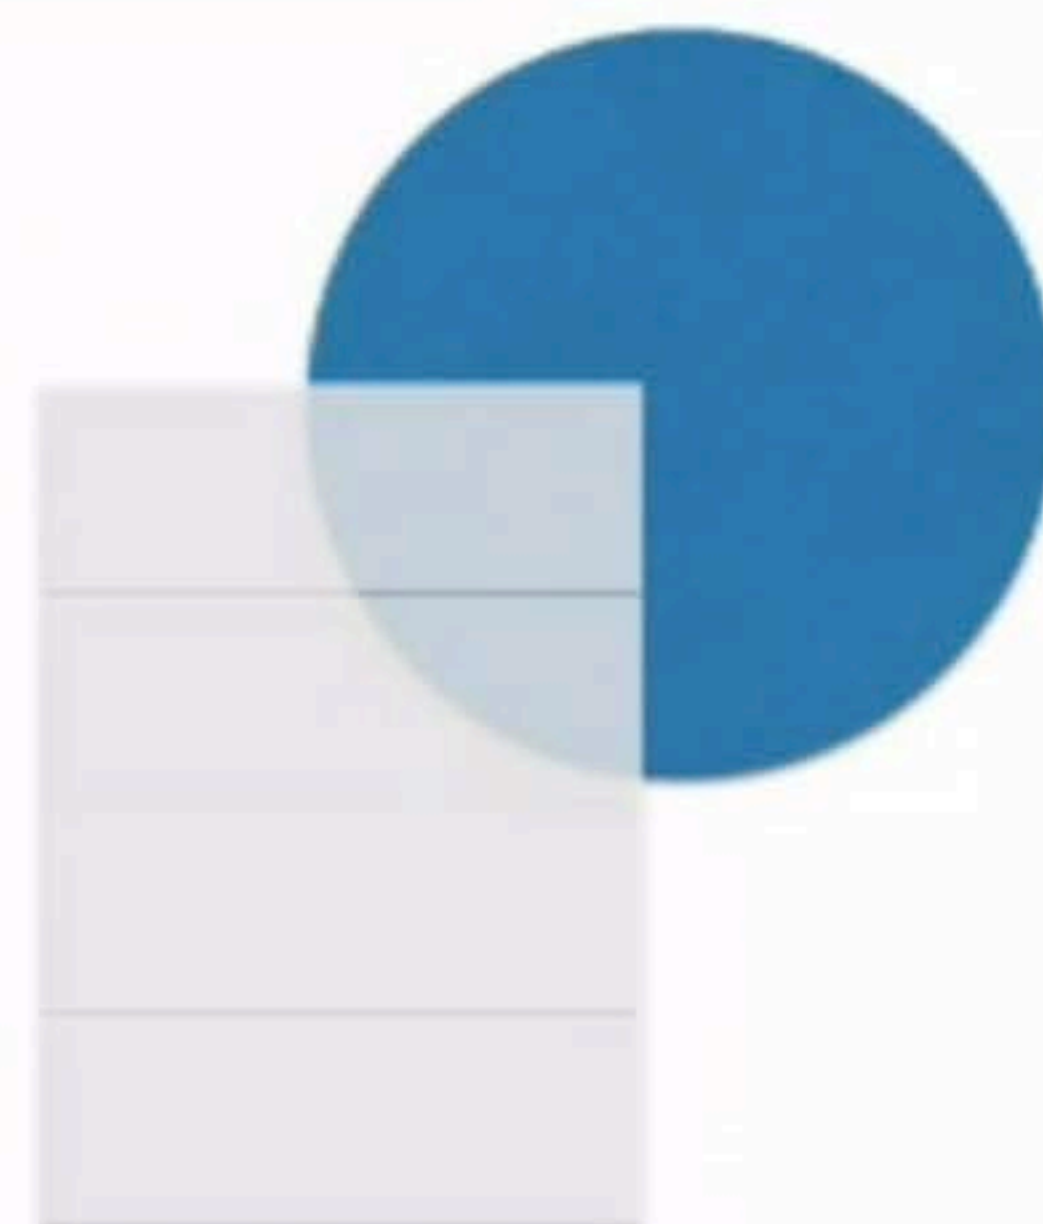
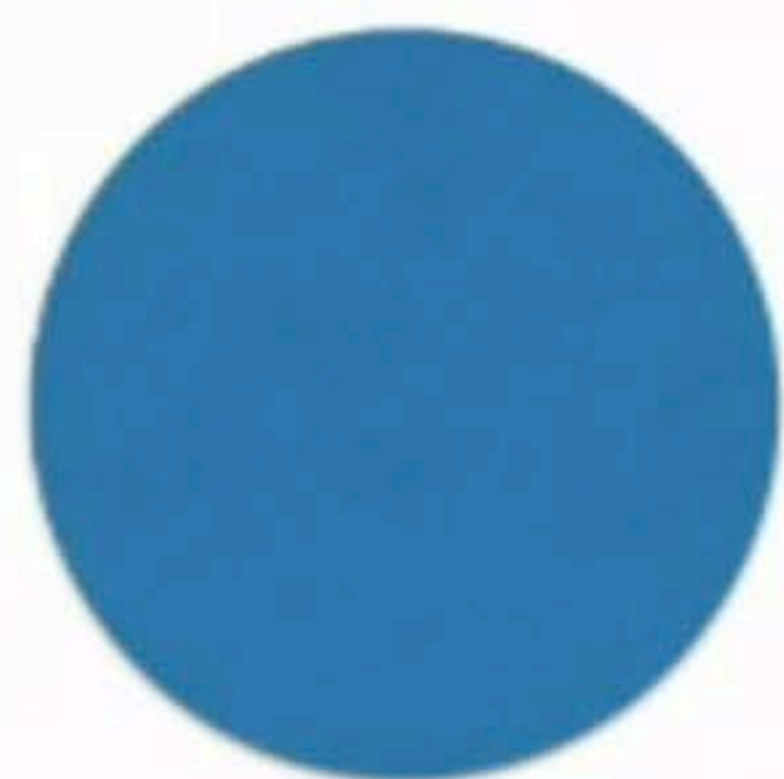
The mailbox in Erlang



Erlang concurrency in a nutshell

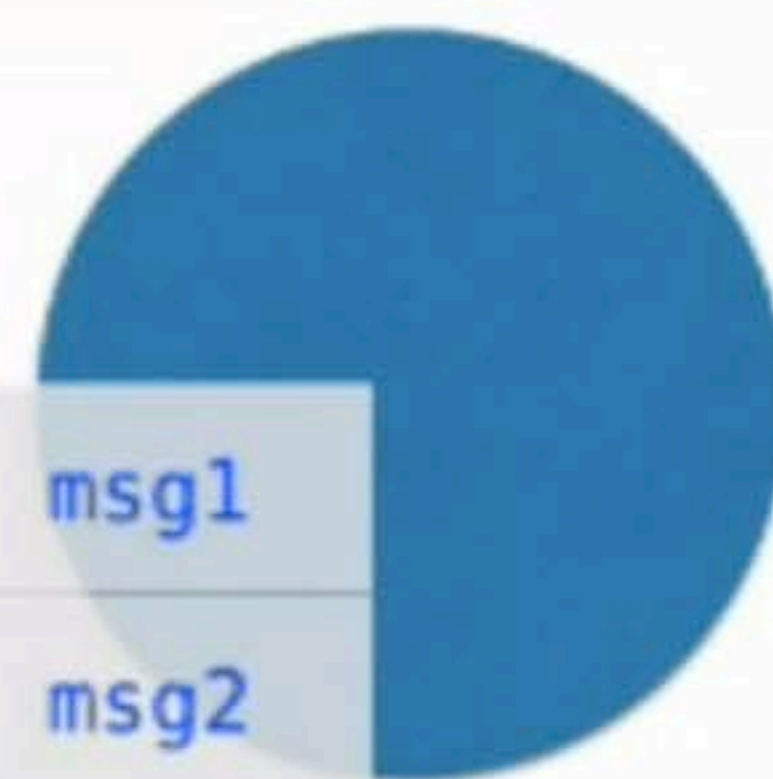
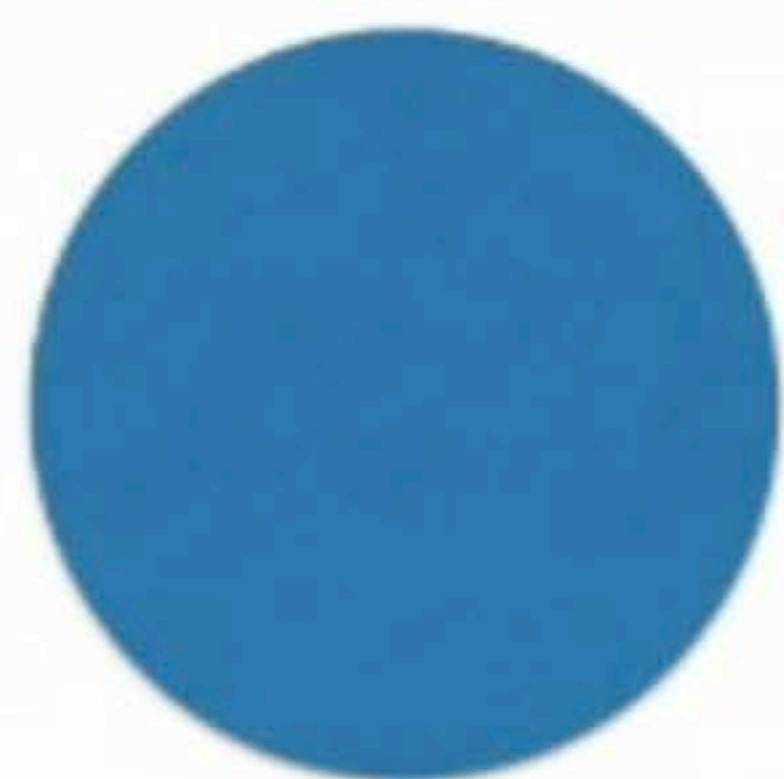
- `spawn` – create a process
- `!` – send a message
- `receive` – process a message
- `self()` – give the Pid of a process





```
Pid ! msg1,  
Pid ! msg2,  
...
```

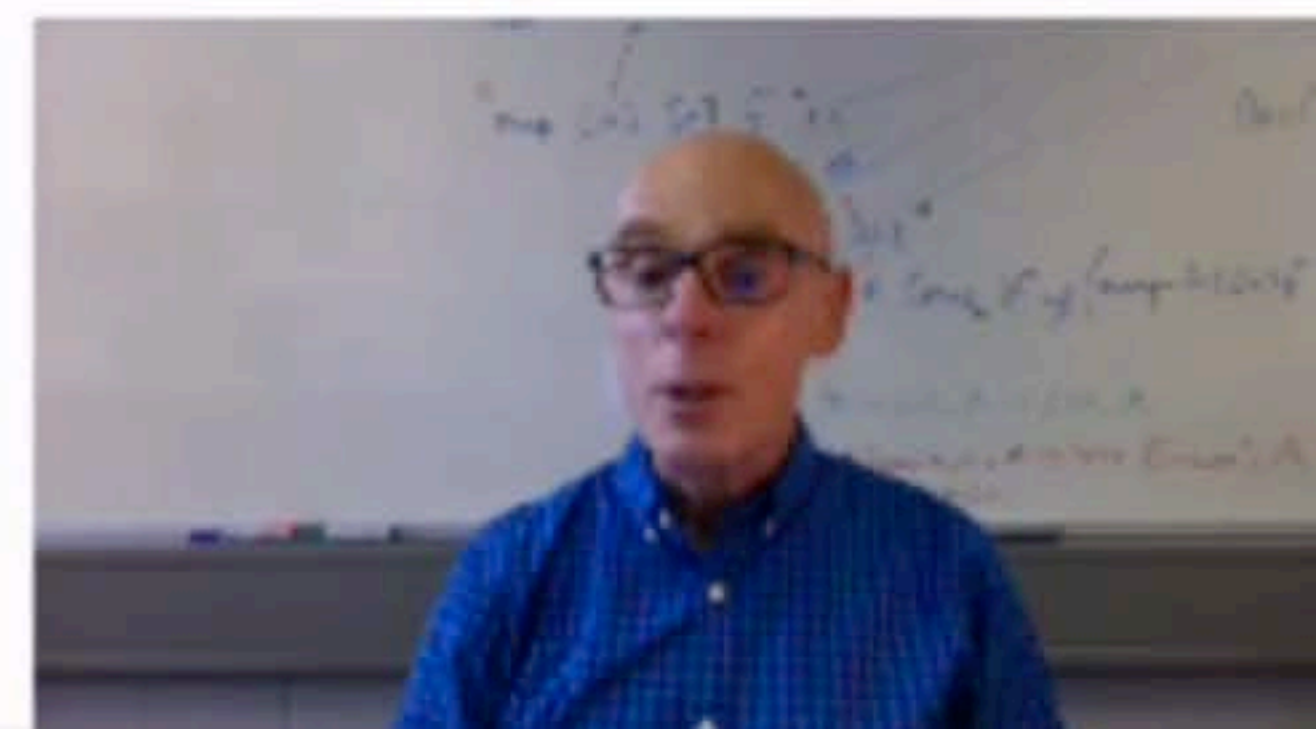


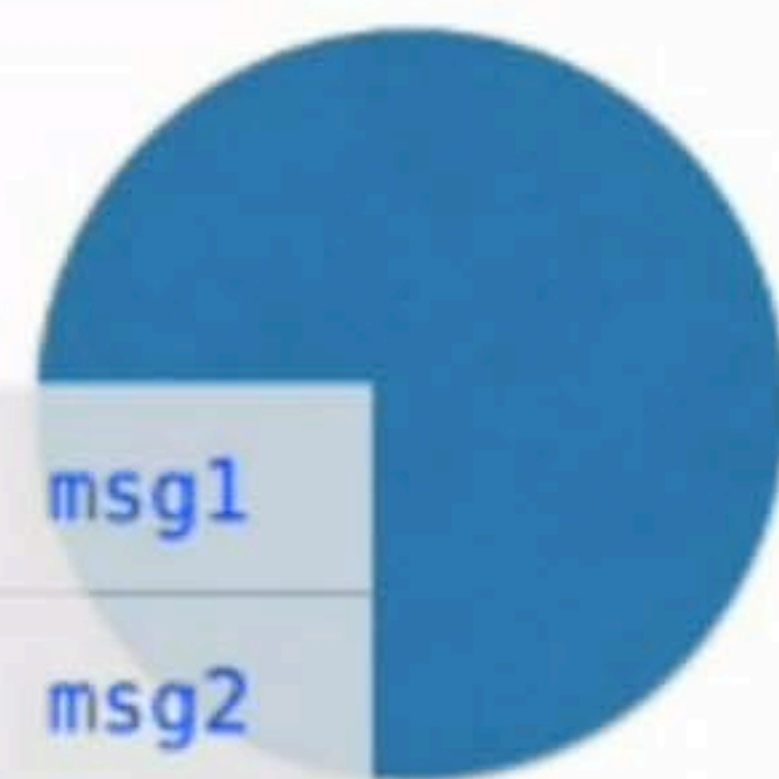


msg1

msg2

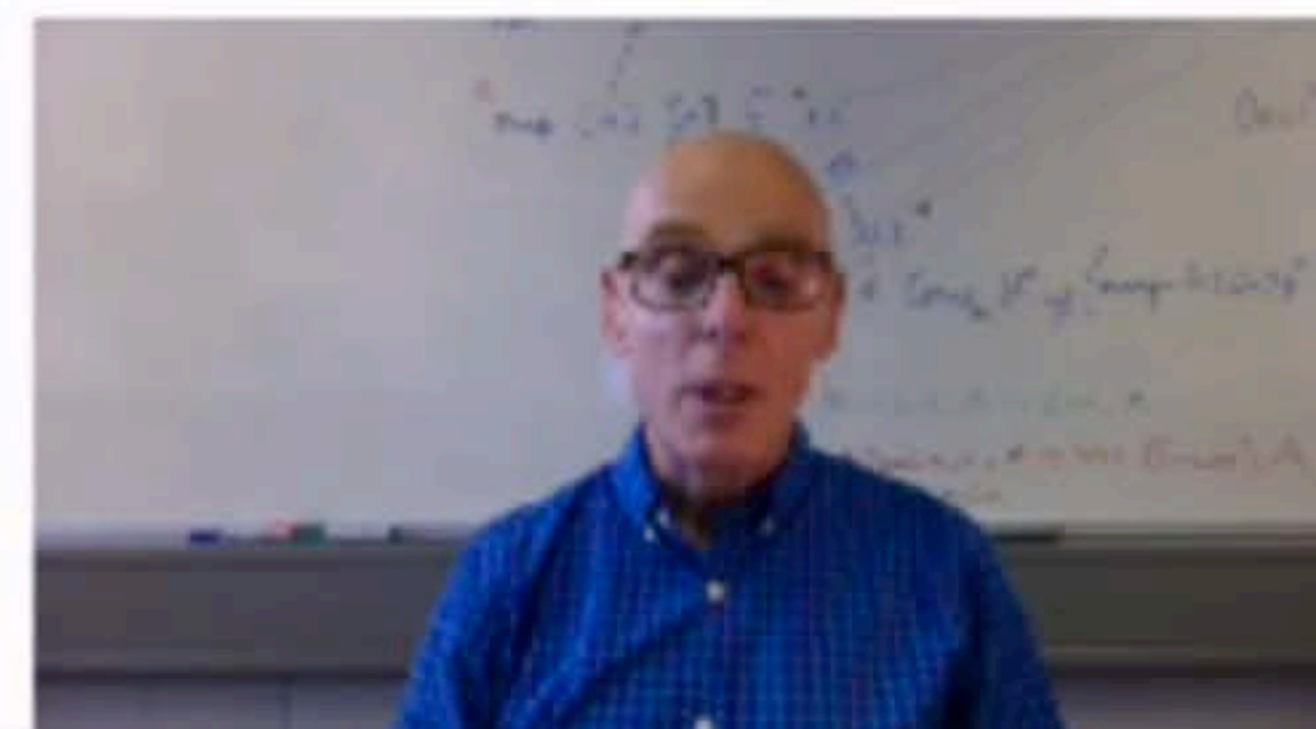
...

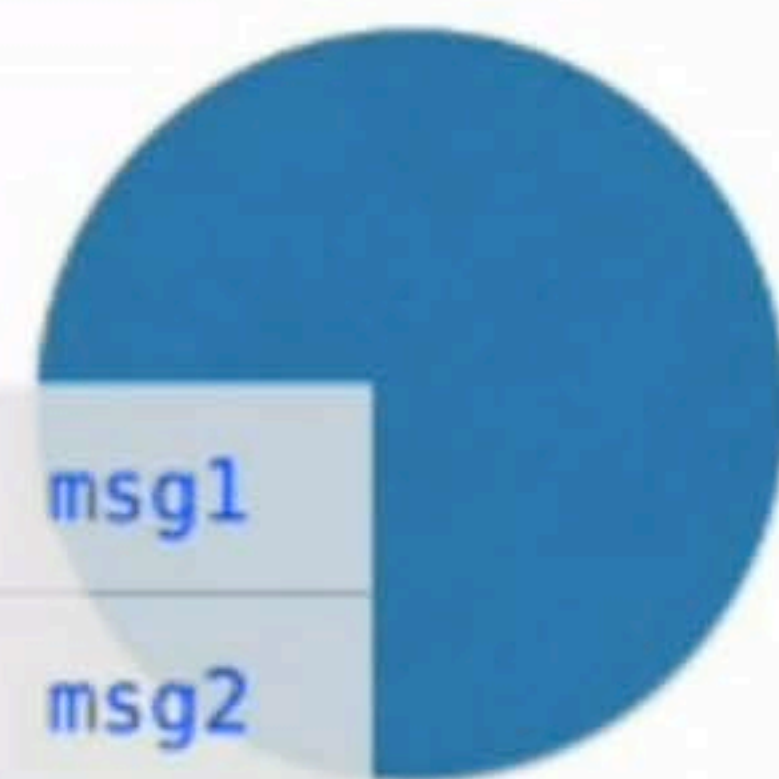
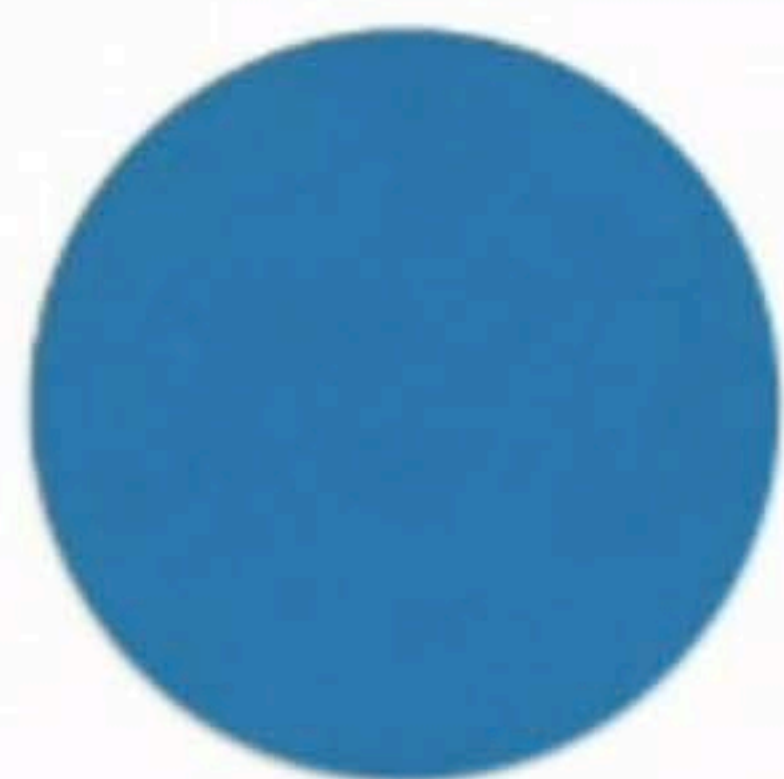




...

```
receive  
  Msg -> ...  
end
```



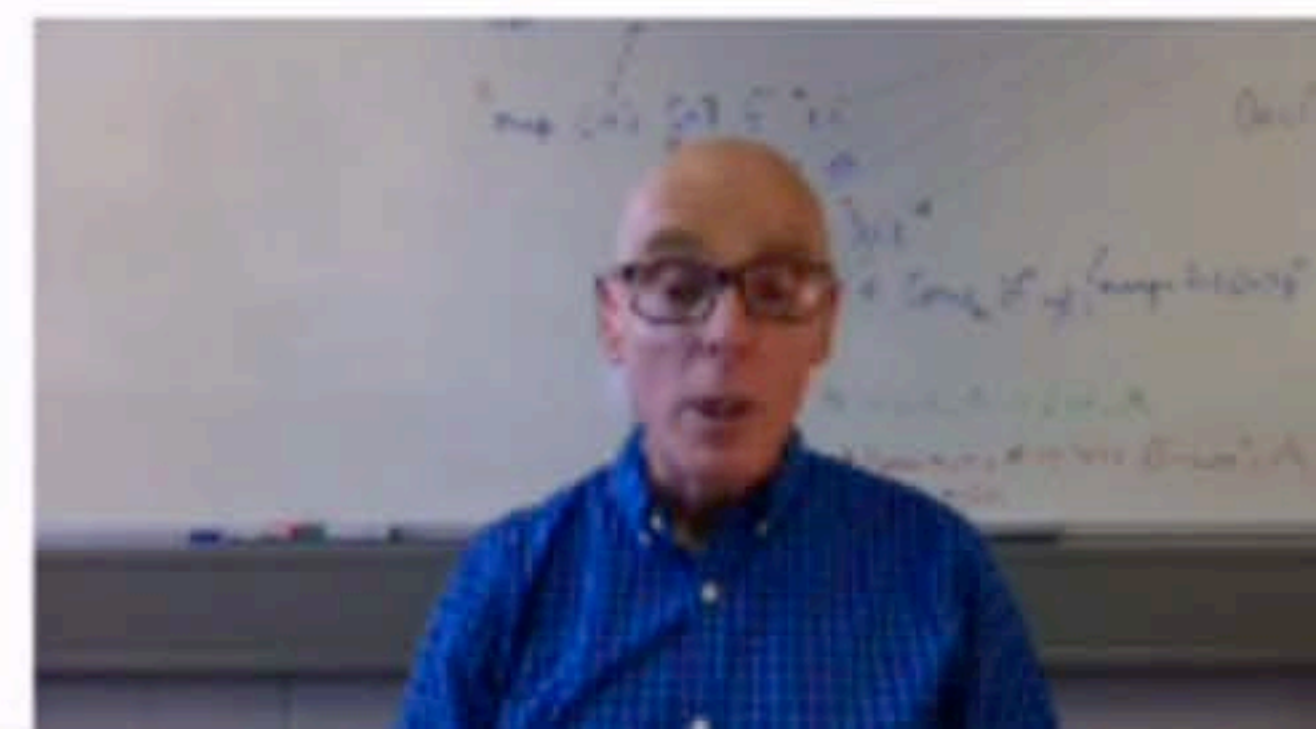


msg1

msg2

...

```
receive  
  msg2 -> ...  
end
```

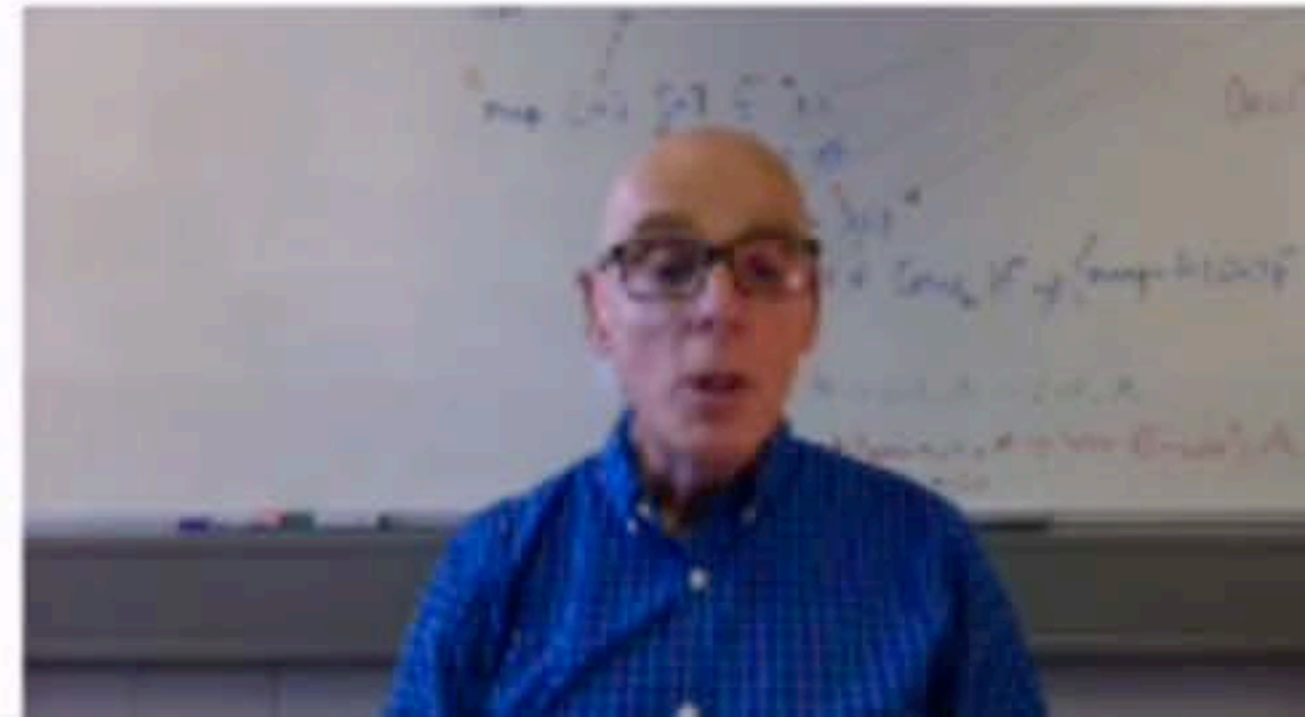


How does message passing work?

When a message is sent to a process, it is added to the mailbox of the process ...

... and messages are stored in the order that they arrive.

Messages in the mailbox are then handled by **receive** statements.



The details of `receive`

Messages *do not* have to be handled in mailbox order.

Try the clauses of the `receive` in turn on the first message ...

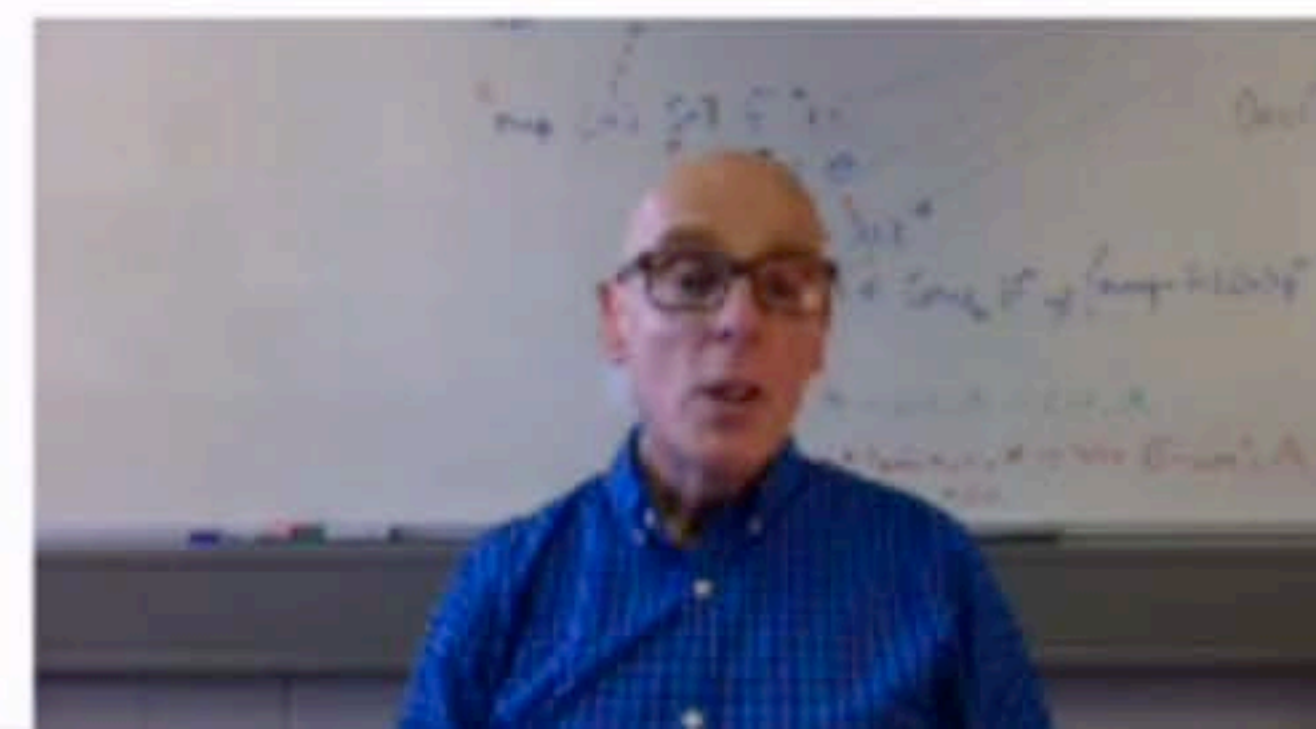
- ... and apply the first matching clause;

- ... if no clause matched, go to the second message and repeat ...

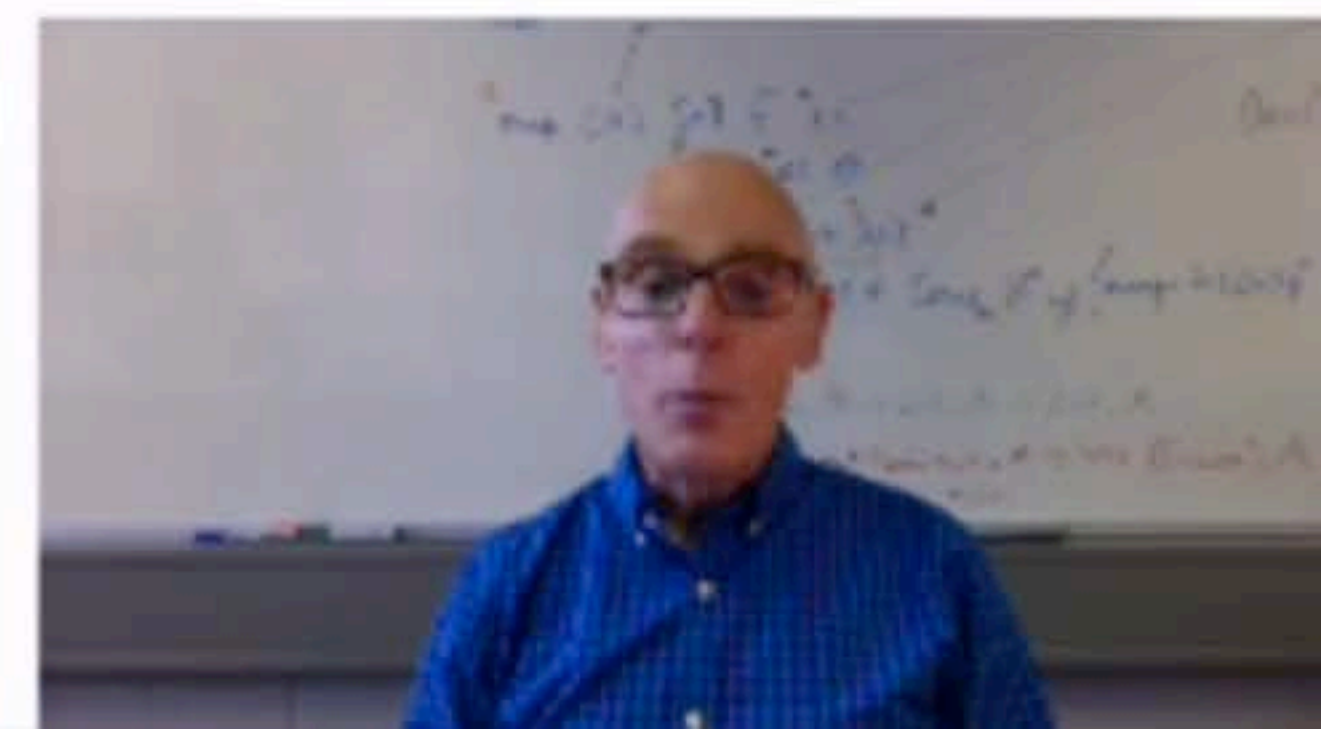
- ... if no message matches any clause, wait for an incoming message.



Scenario 1



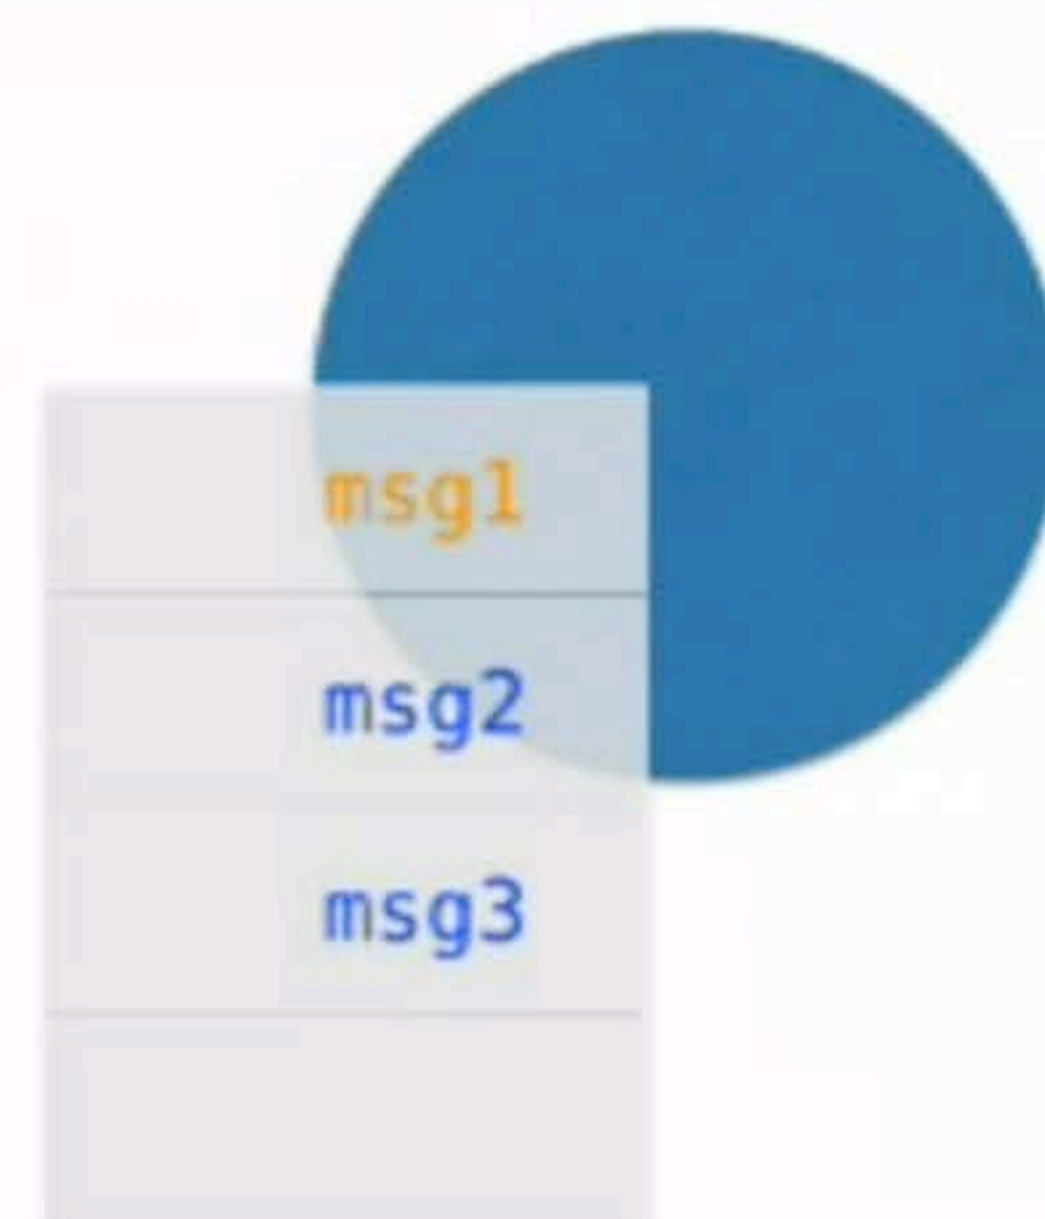
```
receive  
  msg2 -> ...;  
  Msg  -> ...;  
  ...  
end
```




```
receive  
  msg2 -> ...;  
  Msg  -> ...;  
  ...  
end
```




```
receive  
  msg2 -> ...;  
  Msg  -> ...;  
  ...  
end
```

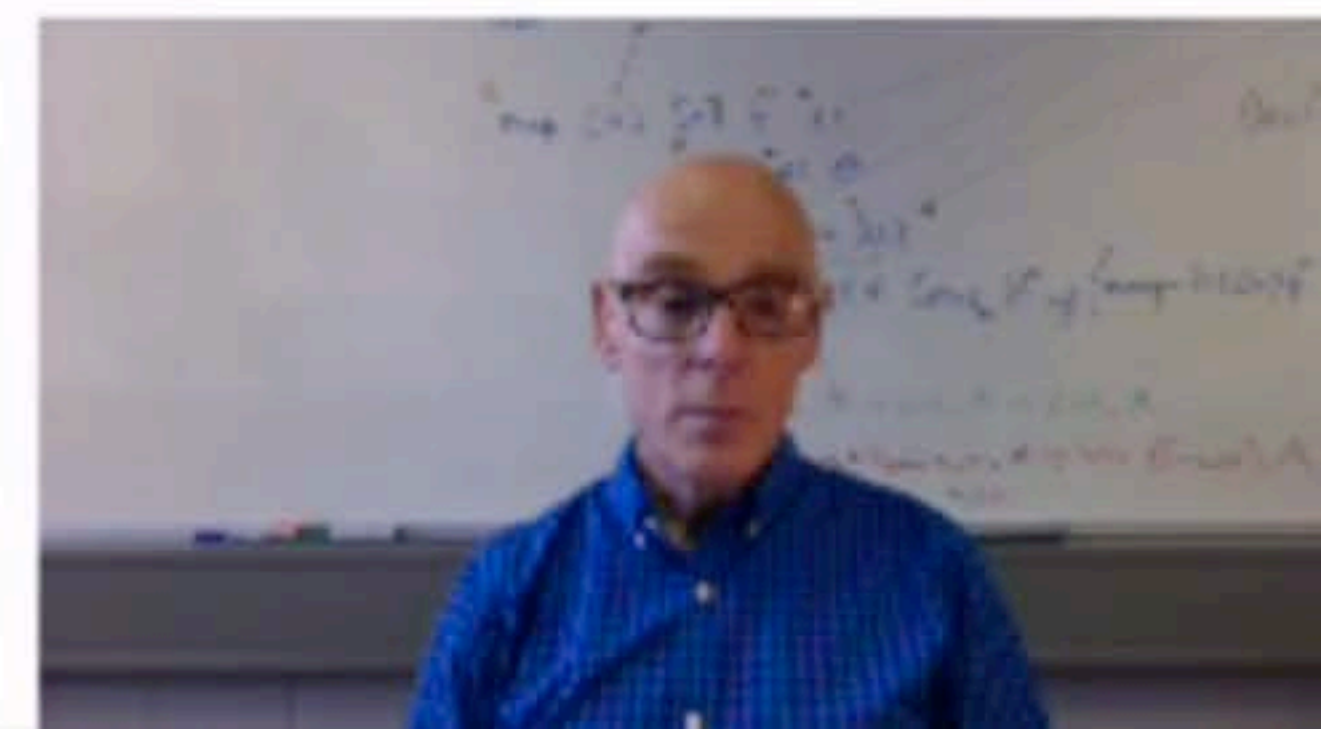


msg2

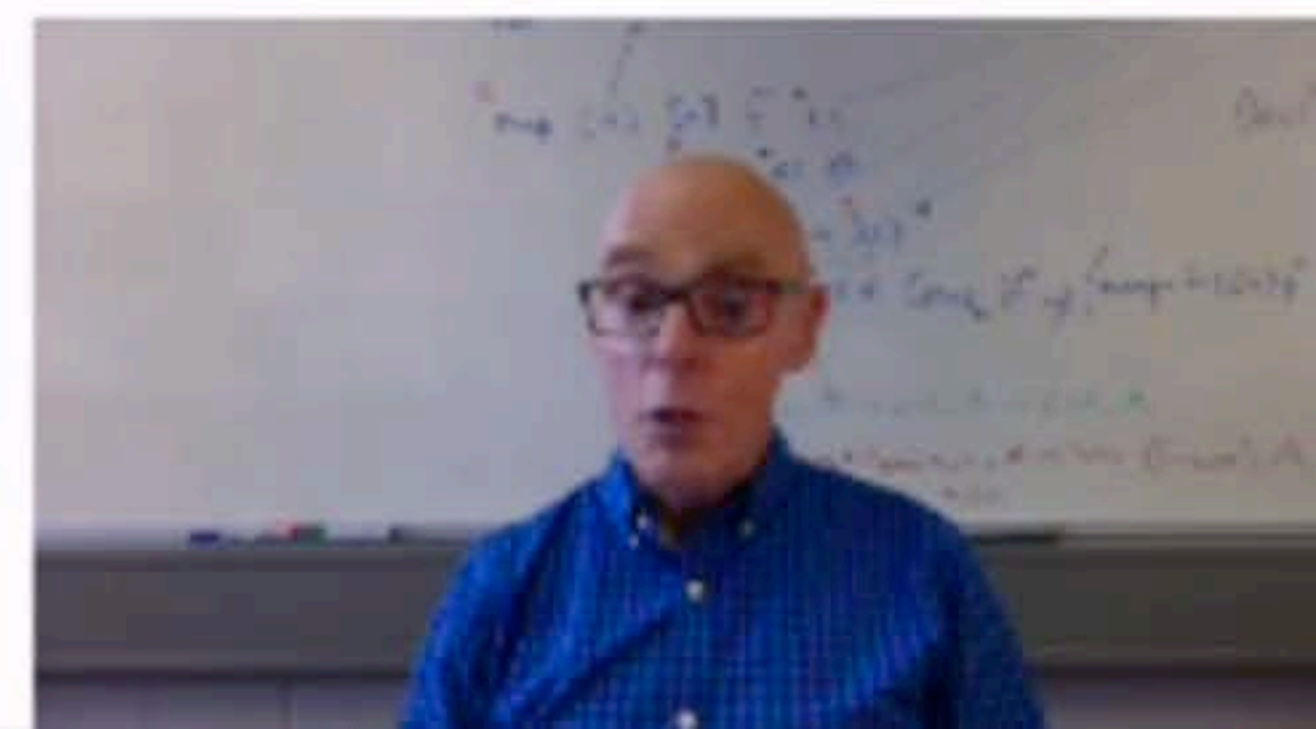
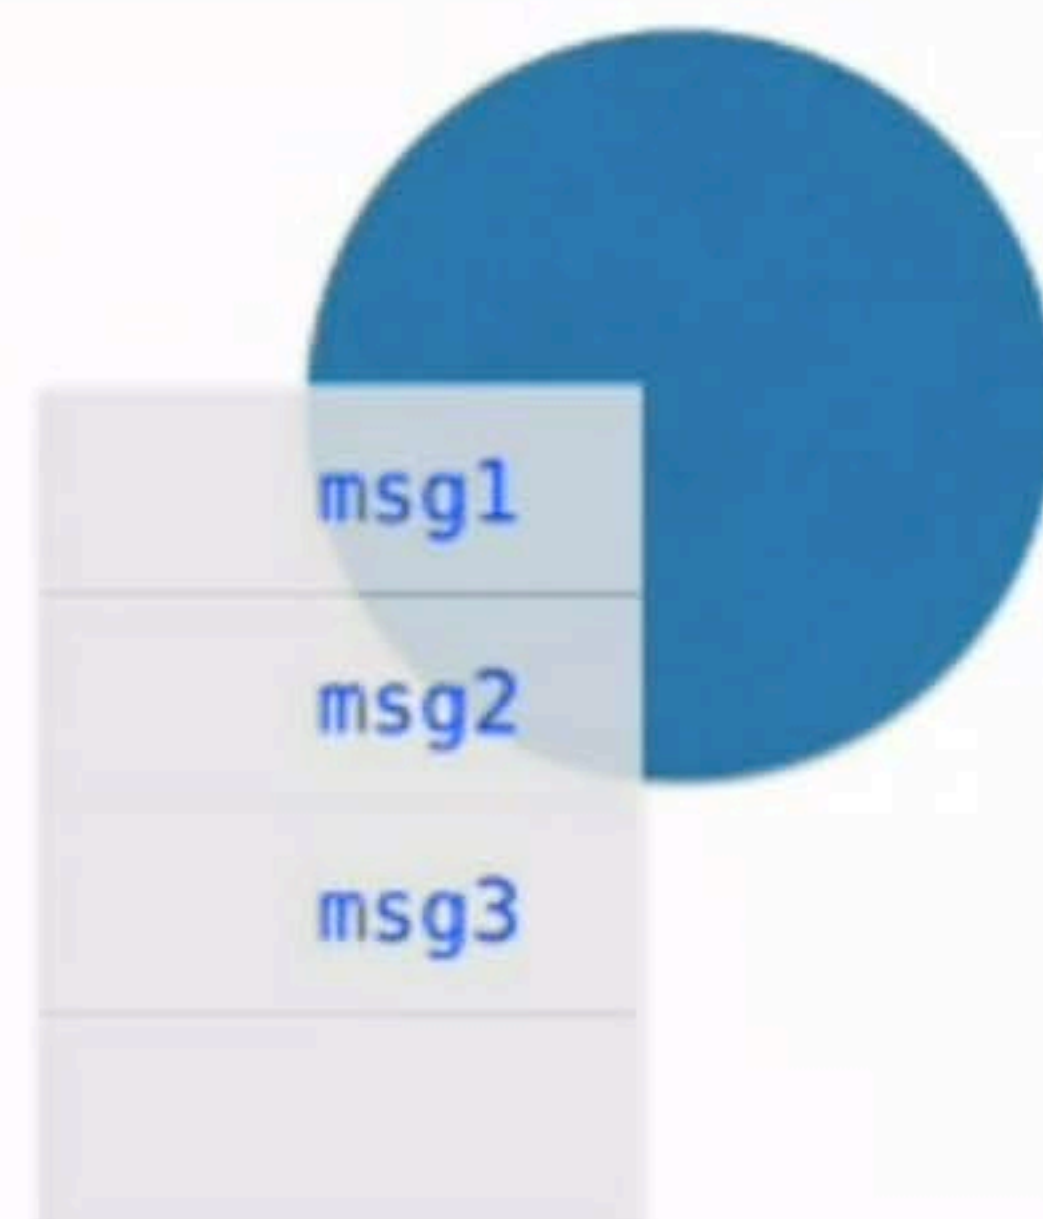
msg3



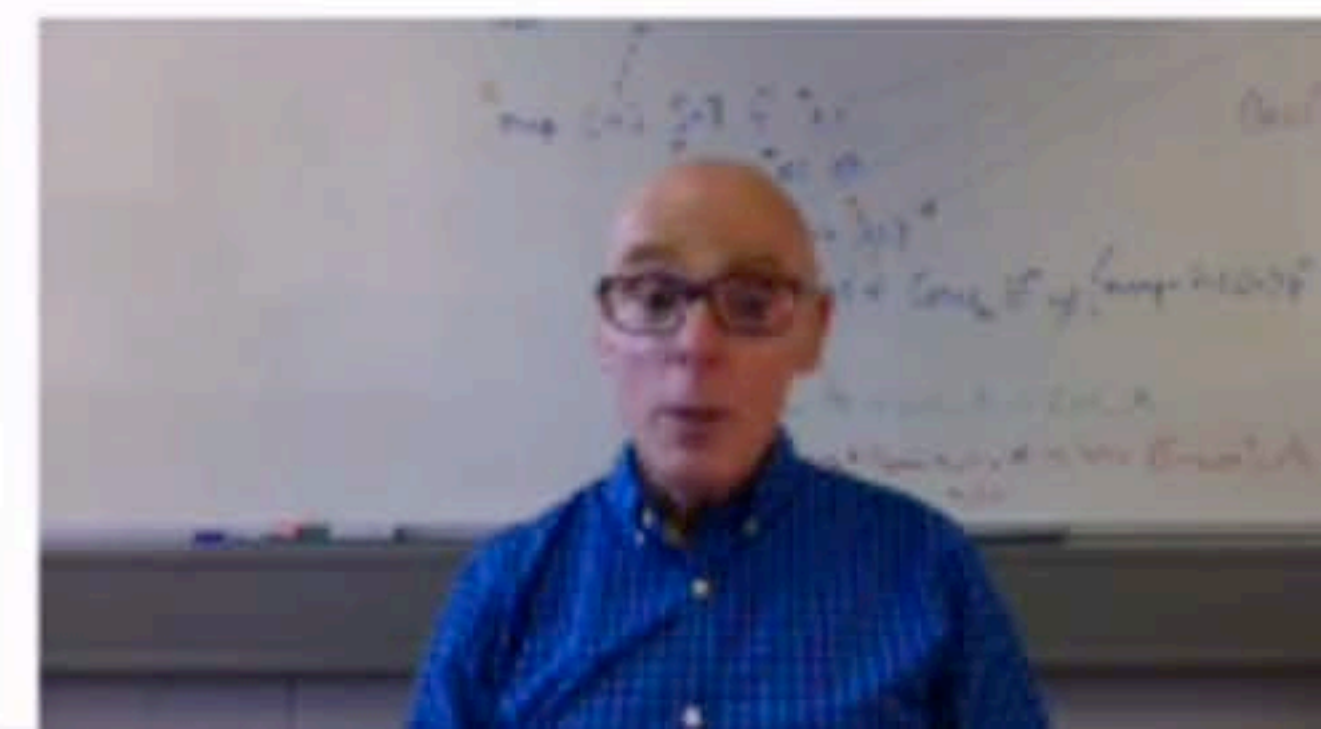
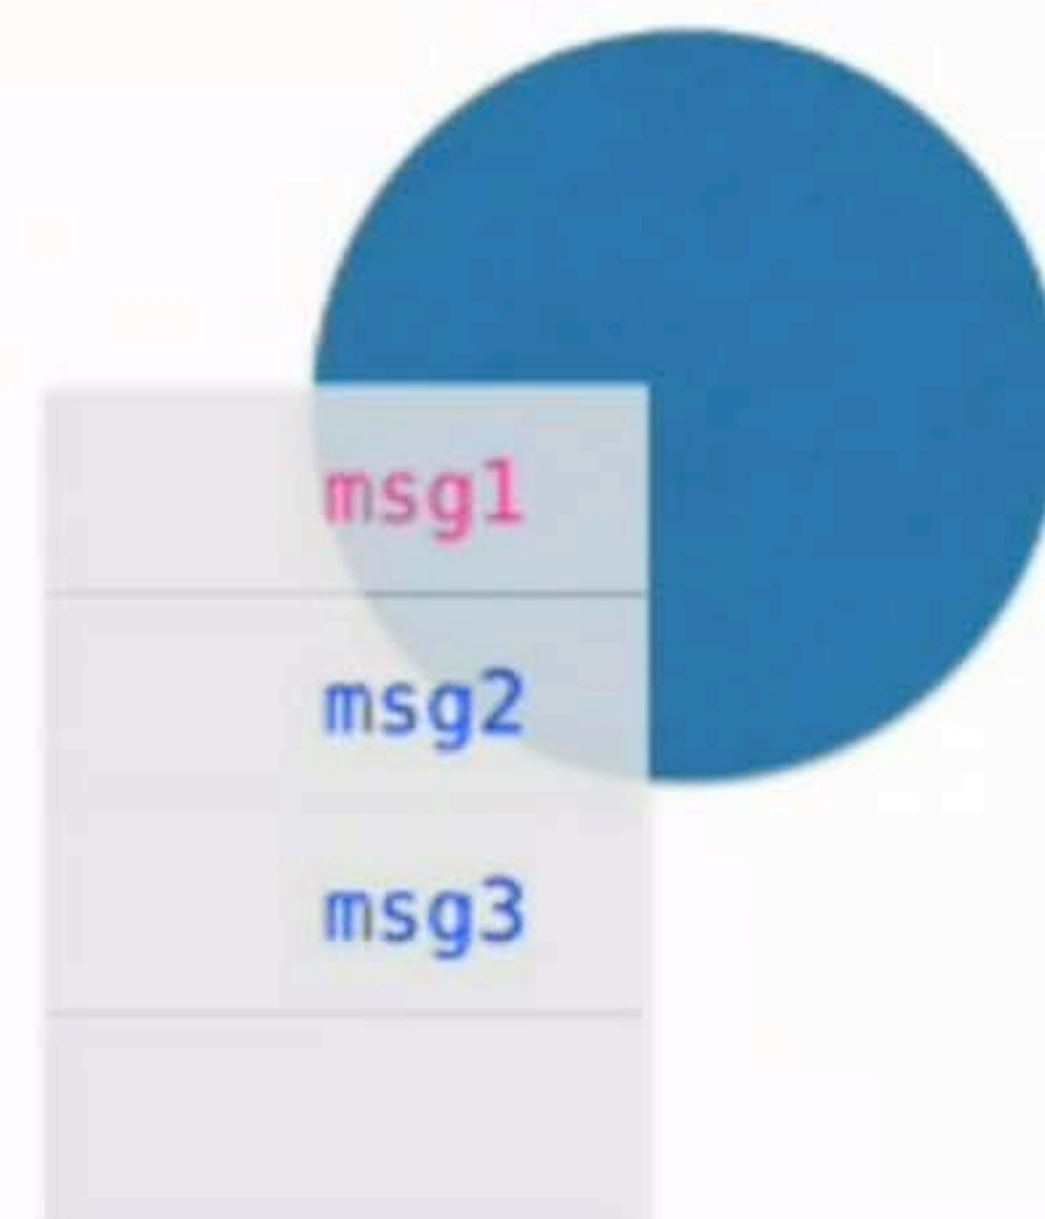
Scenario 2



```
receive  
  msg2 -> ...;  
  msg3 -> ...;  
  msg4 -> ...  
end
```



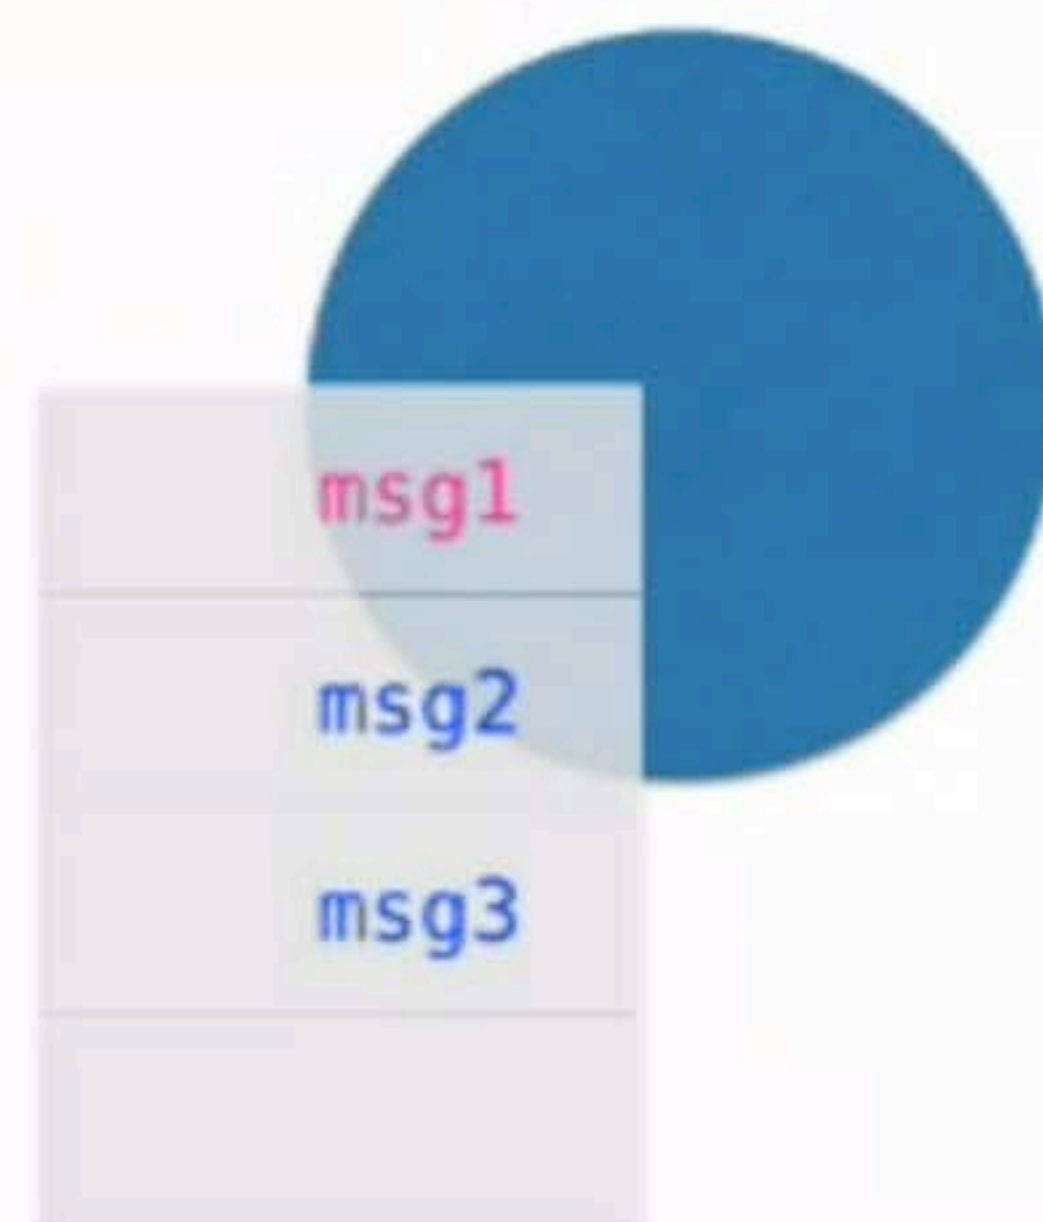

```
receive  
  msg2 -> ...;  
  msg3 -> ...;  
  msg4 -> ...  
end
```



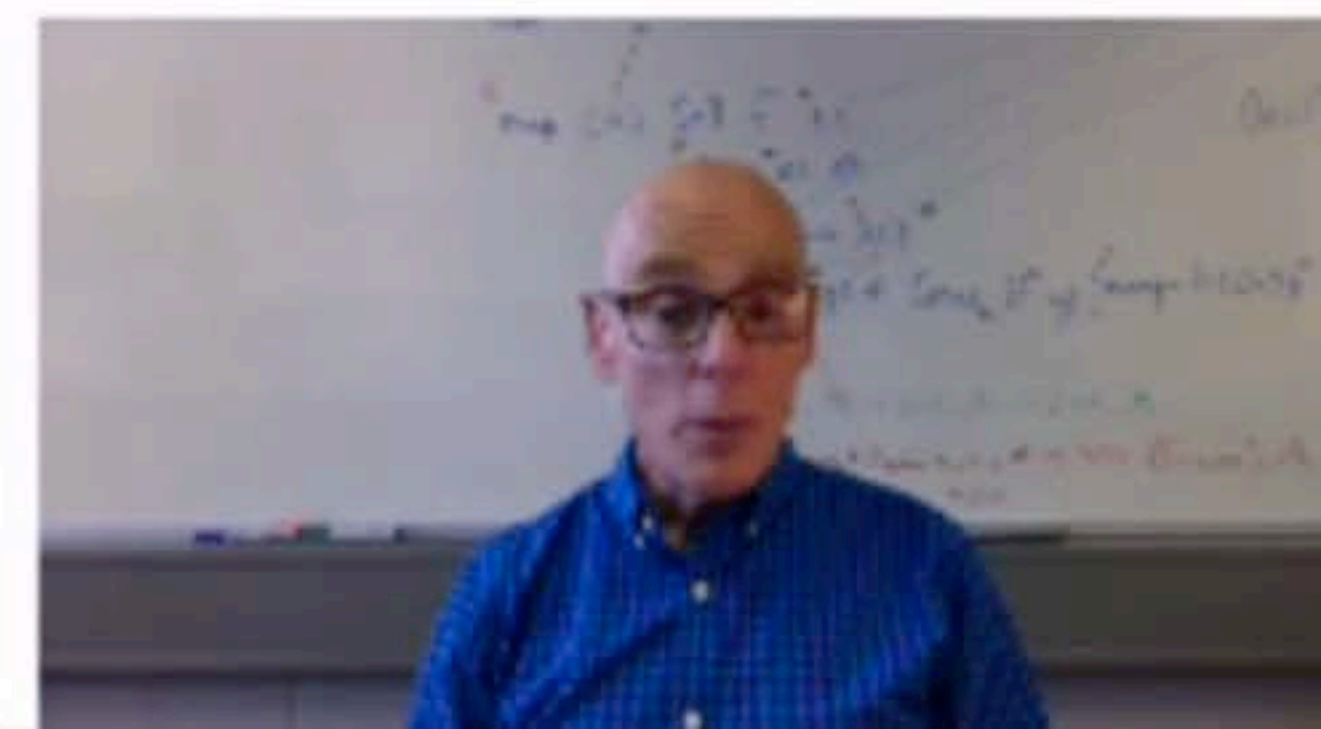
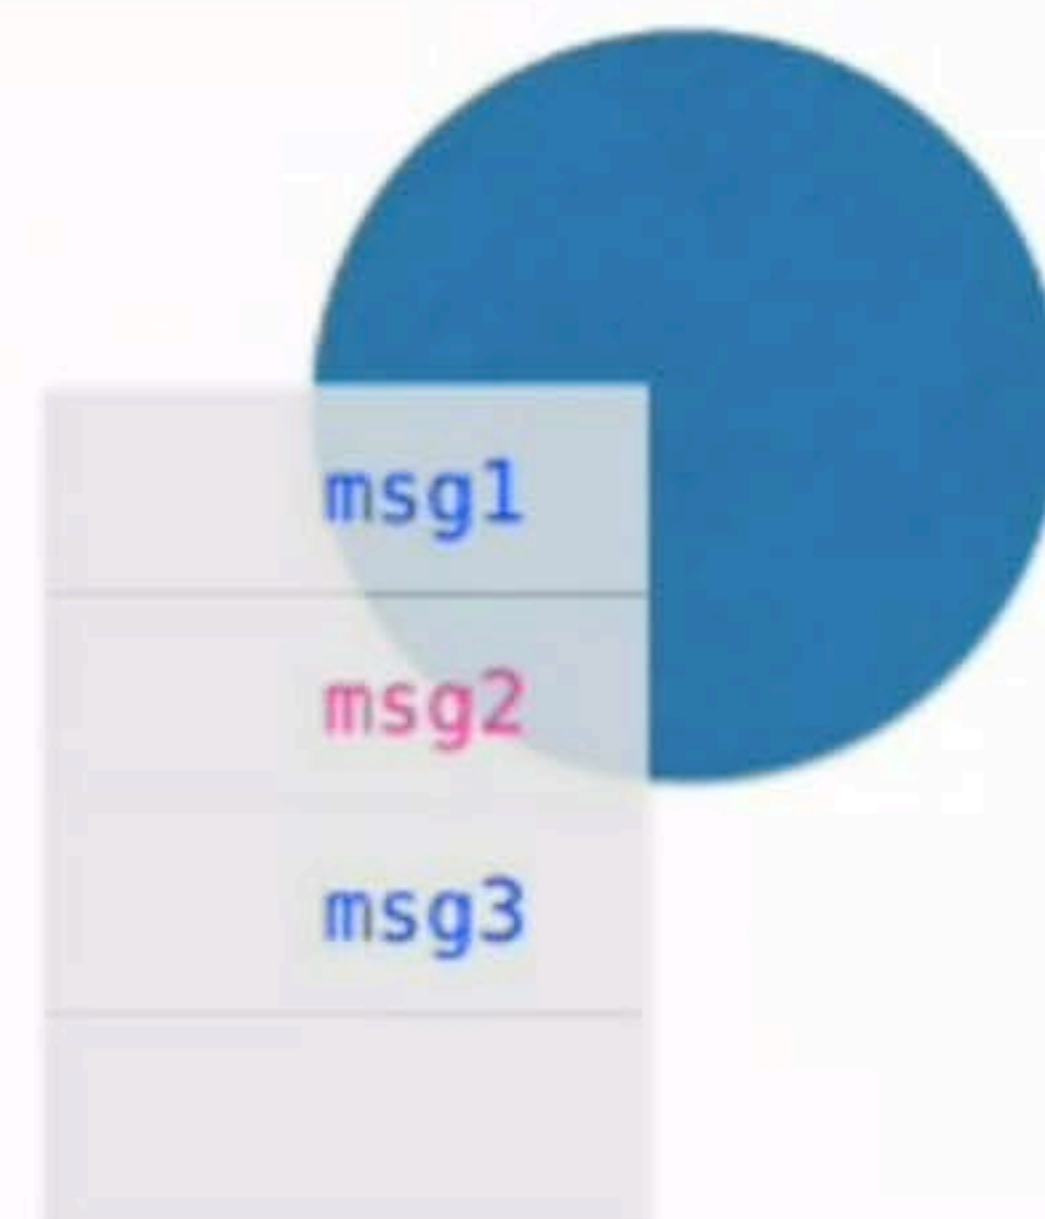
```
receive  
  msg2 -> ...;  
  msg3 -> ...;  
  msg4 -> ...  
end
```



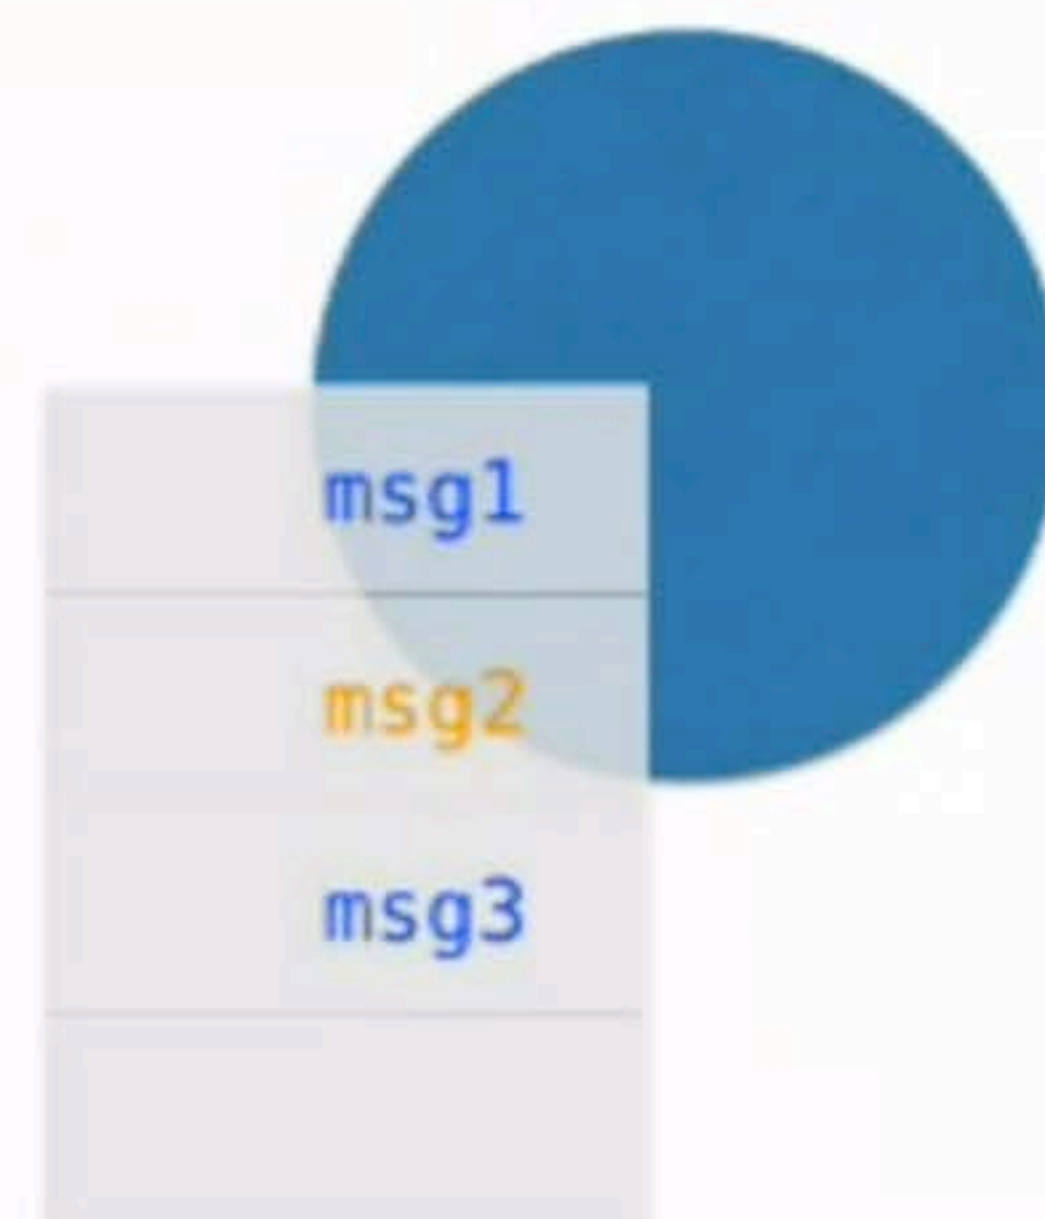
```
receive  
  msg2 -> ...;  
  msg3 -> ...;  
  msg4 -> ...  
end
```



```
receive  
  msg2 -> ...;  
  msg3 -> ...;  
  msg4 -> ...  
end
```




```
receive  
  msg2 -> ...;  
  msg3 -> ...;  
  msg4 -> ...  
end
```



msg1

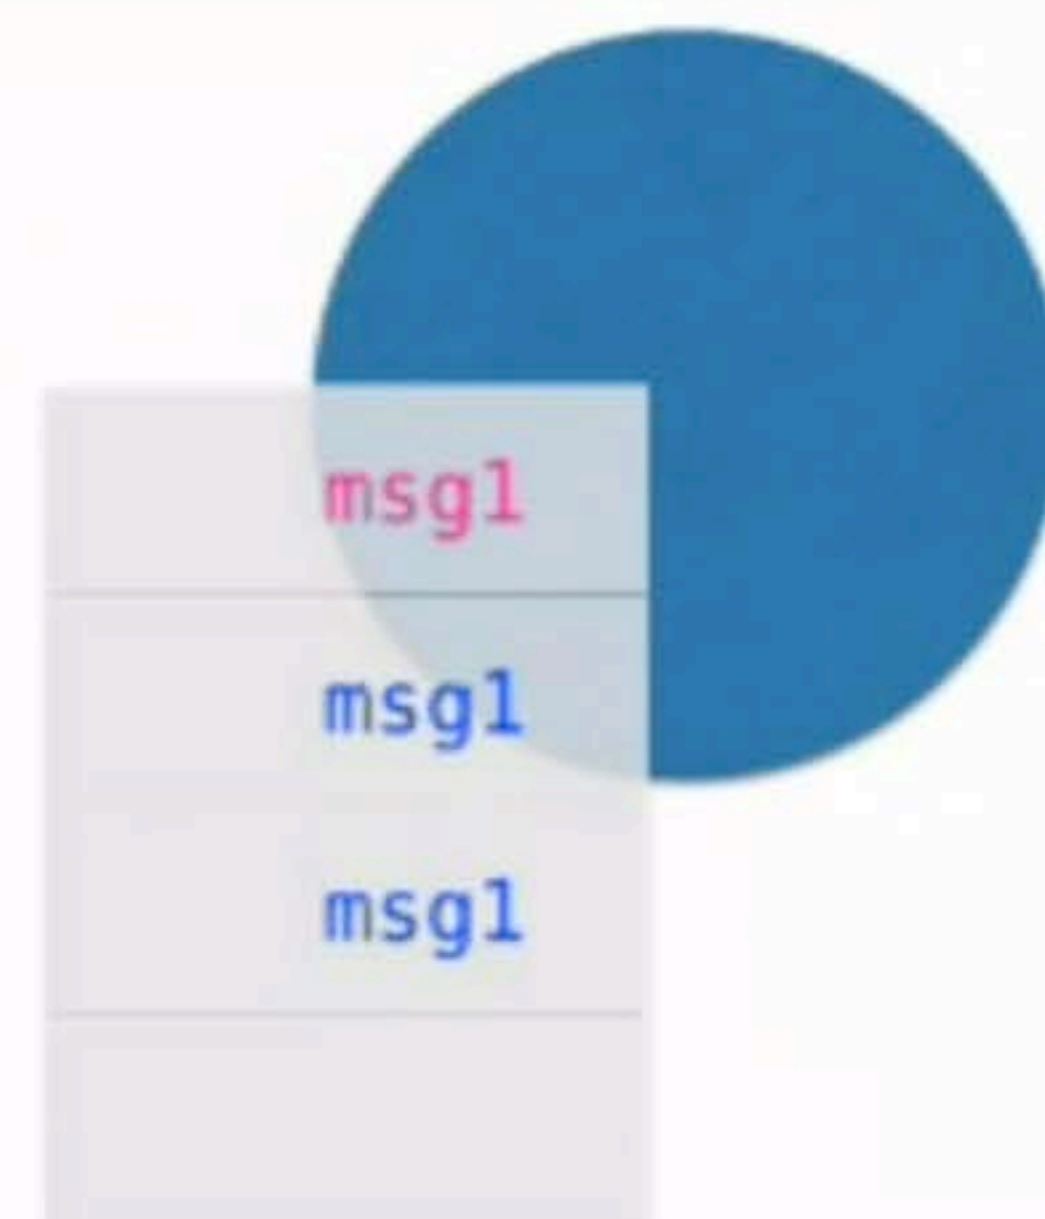
msg3



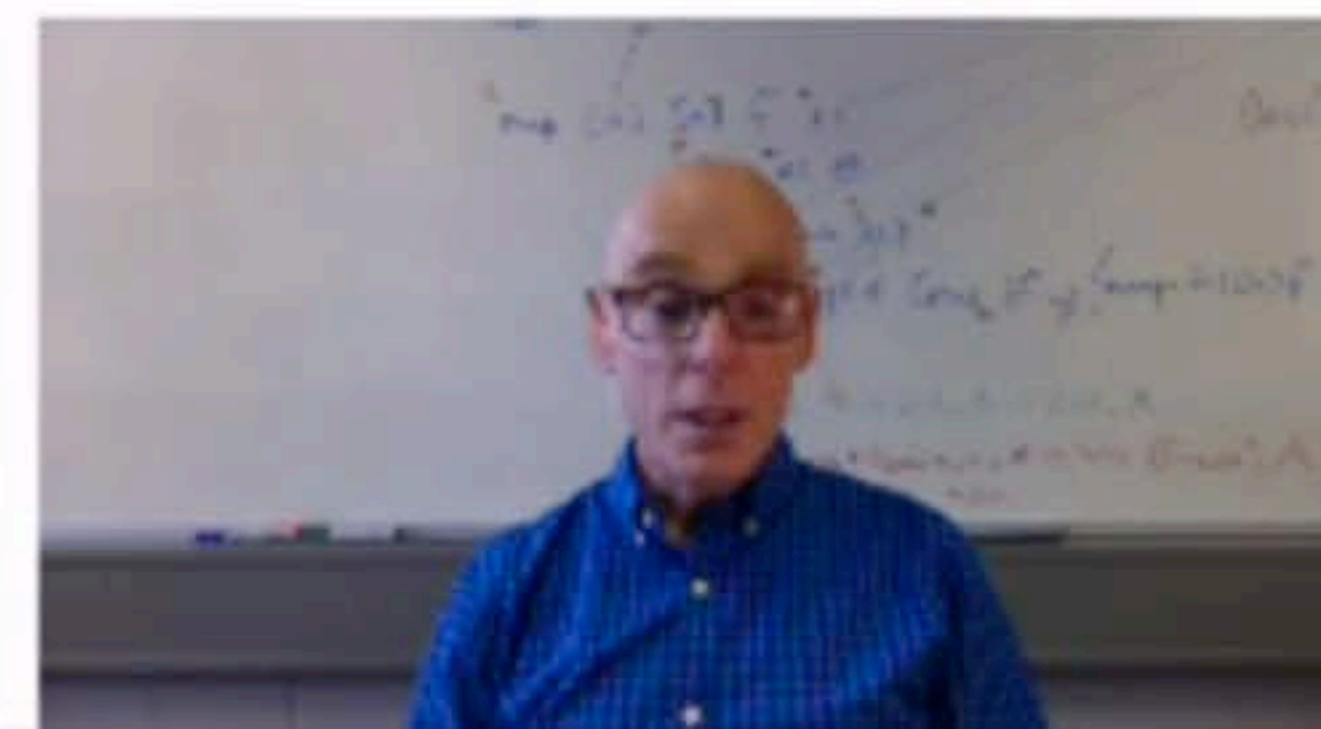
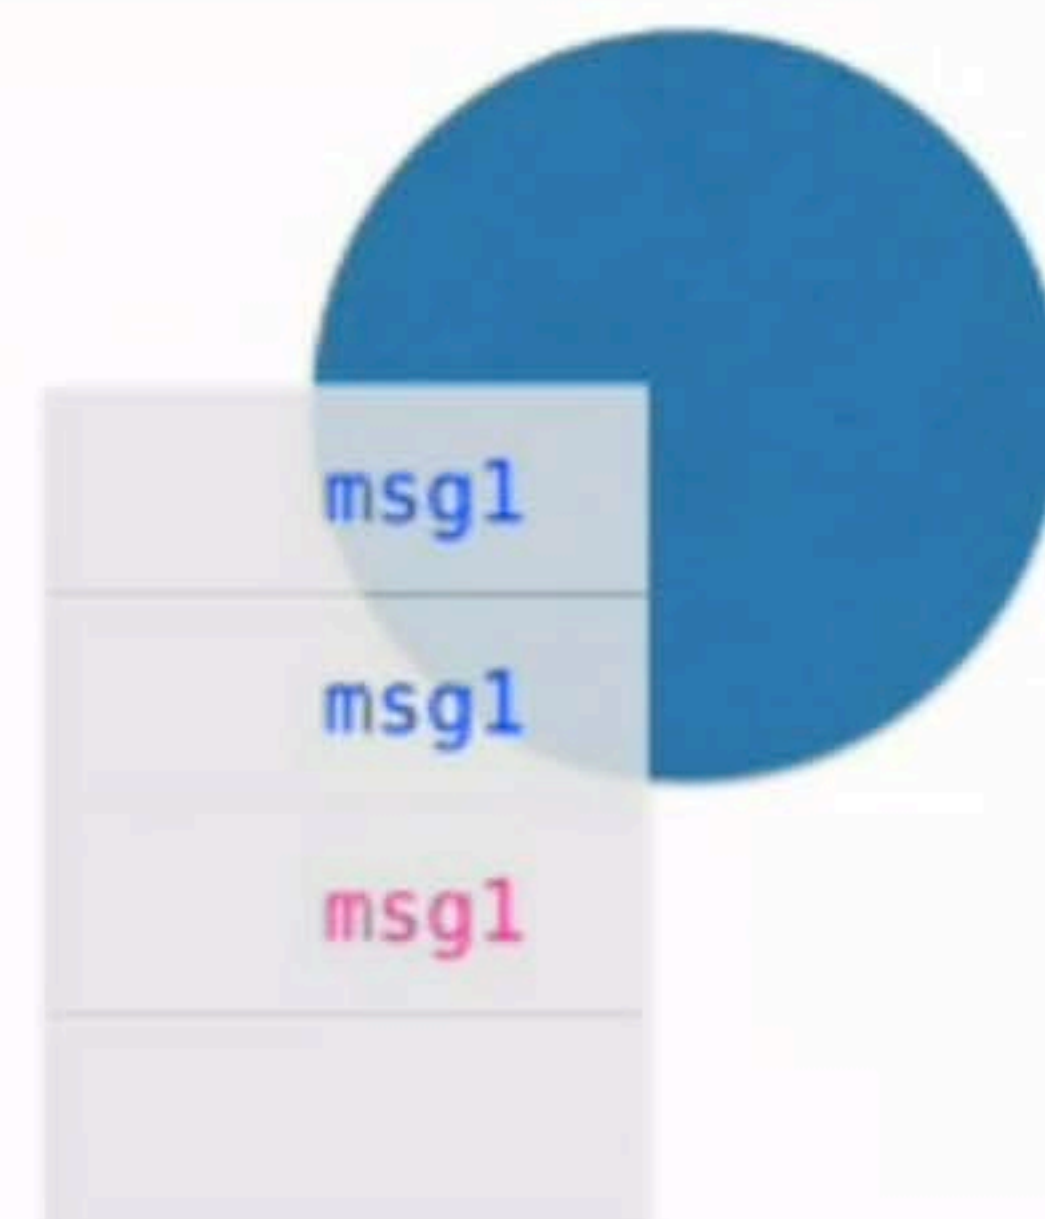
Scenario 3



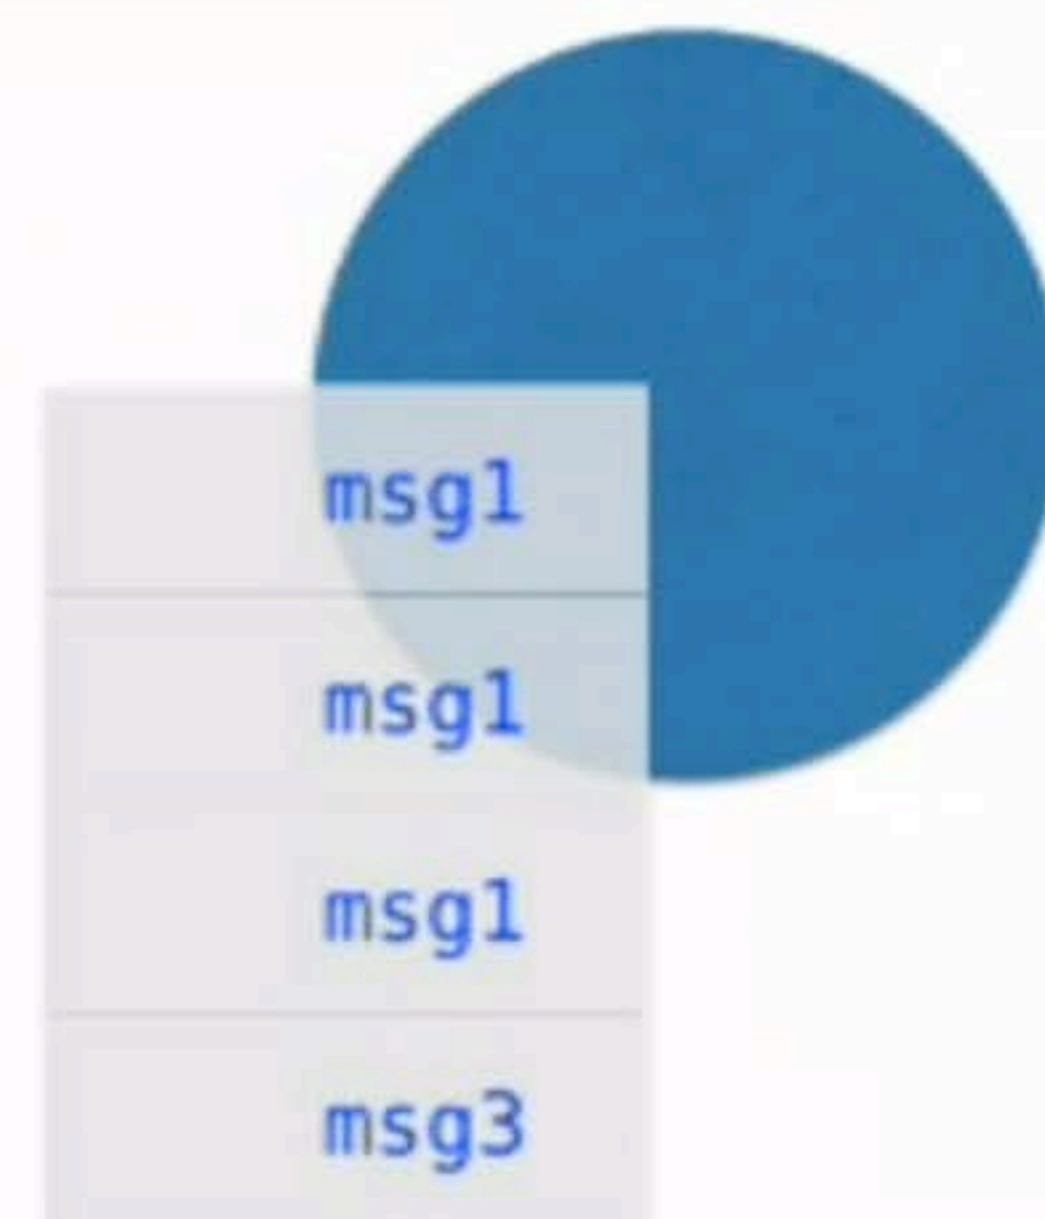
```
receive  
  msg2 -> ...;  
  msg3 -> ...;  
  msg4 -> ...  
end
```



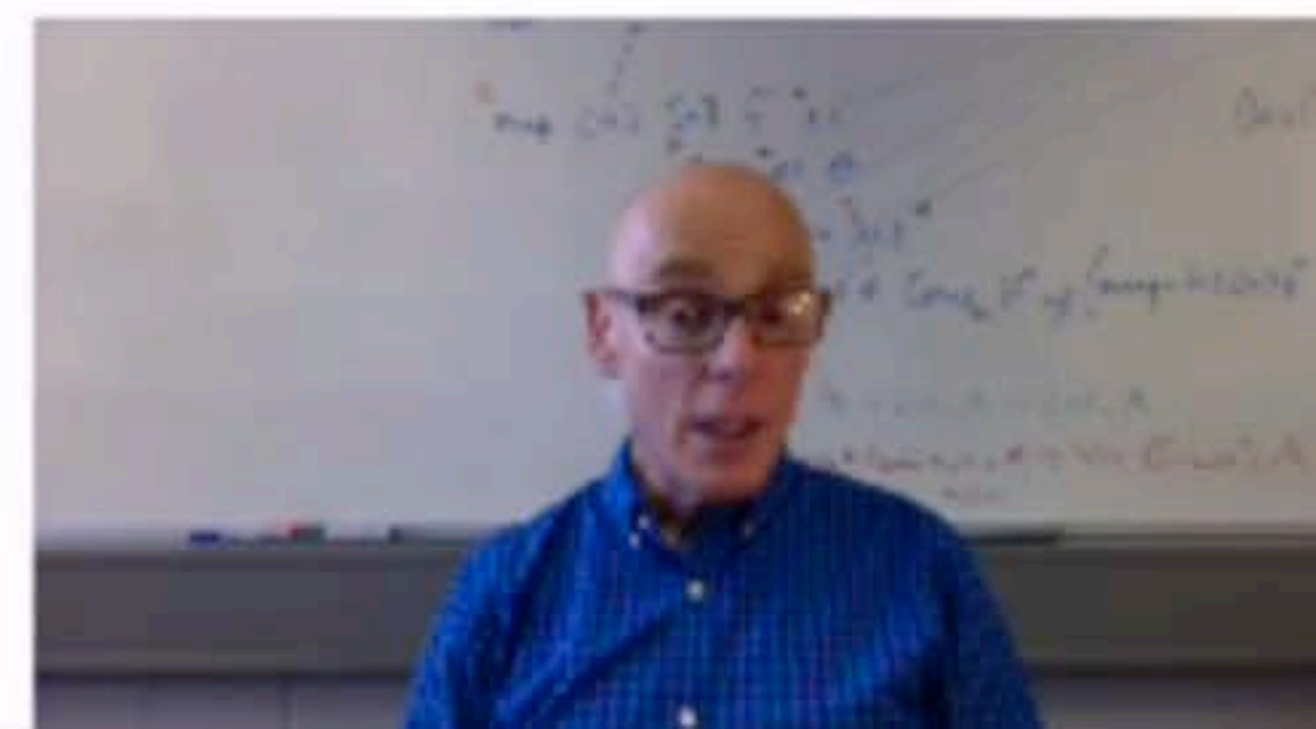
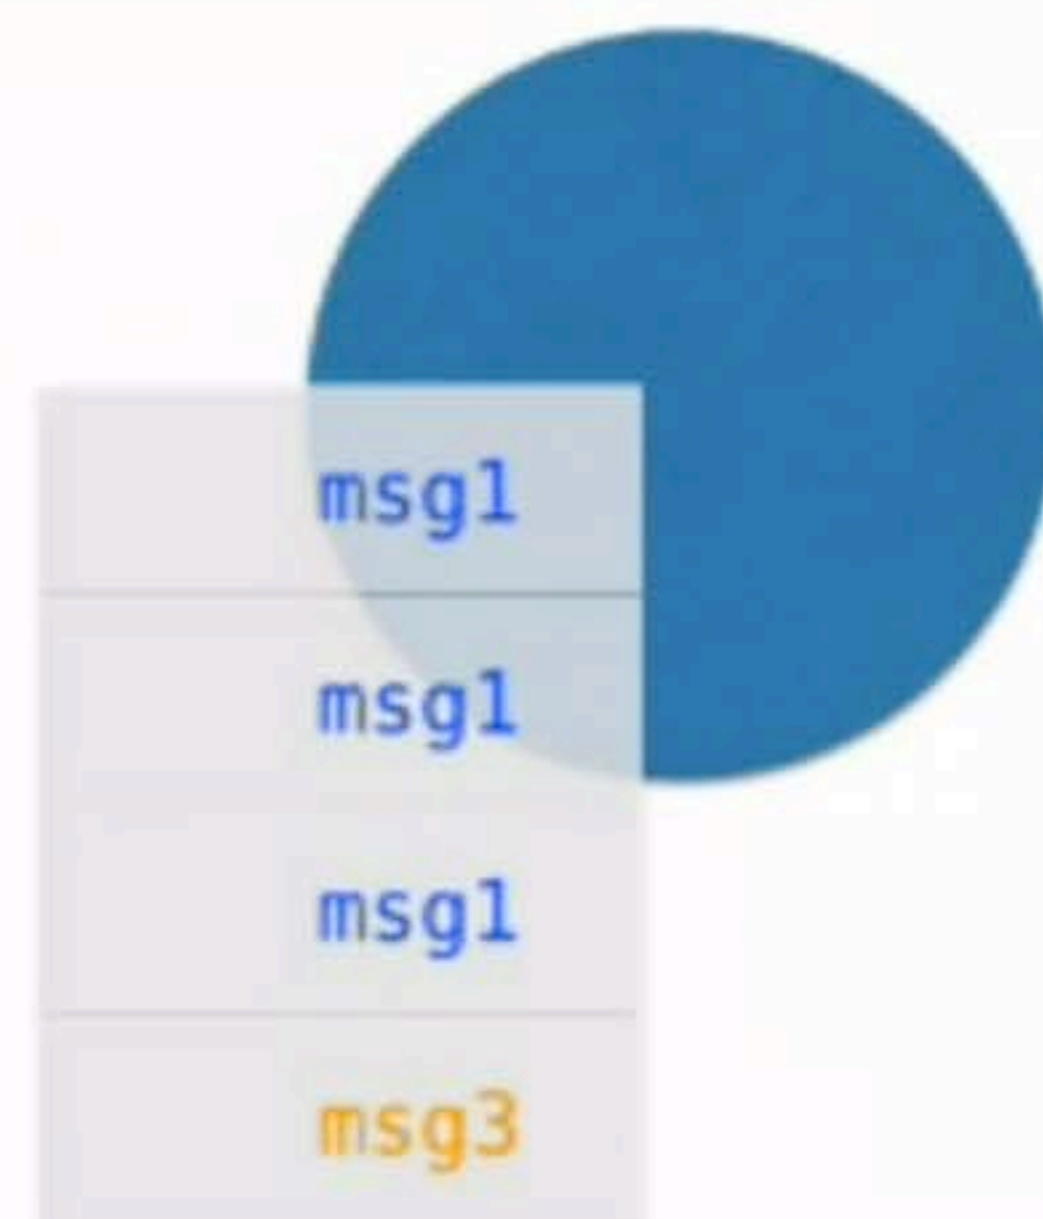

```
receive  
  msg2 -> ...;  
  msg3 -> ...;  
  msg4 -> ...  
end
```



```
receive  
  msg2 -> ...;  
  msg3 -> ...;  
  msg4 -> ...  
end
```



```
receive  
  msg2 -> ...;  
  msg3 -> ...;  
  msg4 -> ...  
end
```



msg1

msg1

msg1



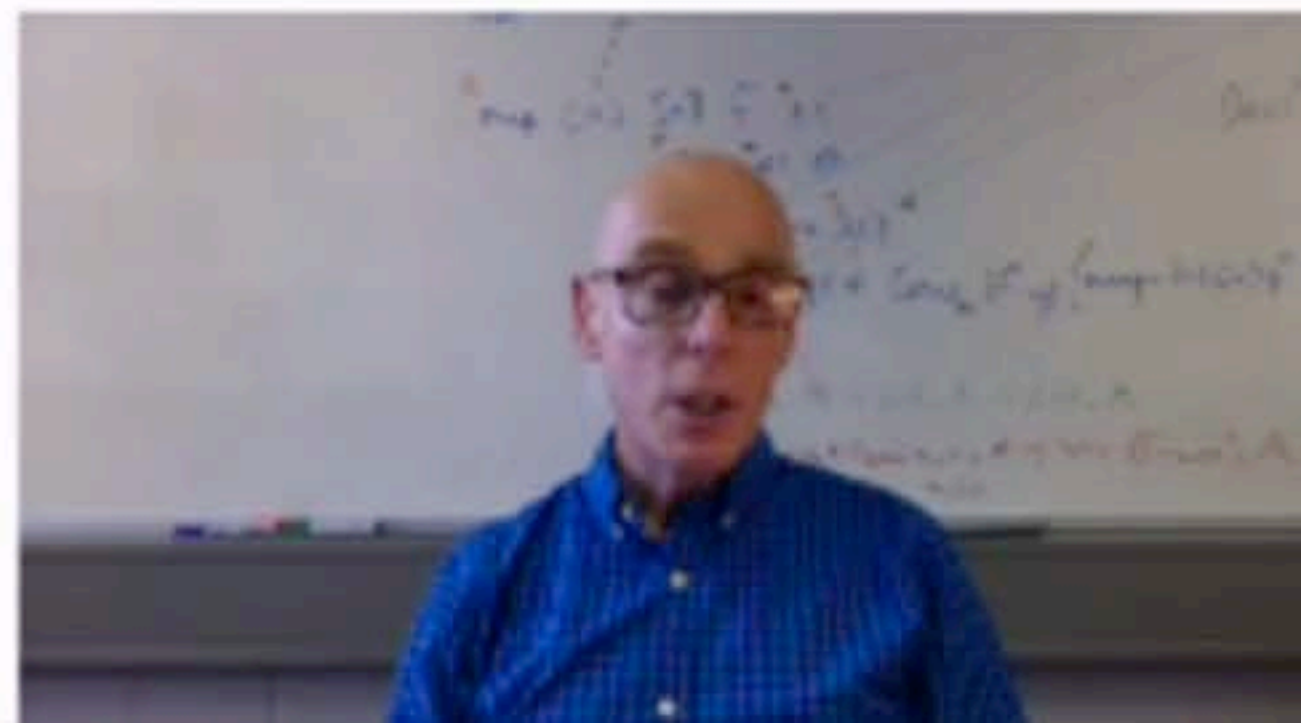
Handling the mailbox

Message sending is asynchronous ...

...and decoupled from message handling using `receive`

Message `receive` is *selective* ...

... allowing messages to be handled in the order that we choose.



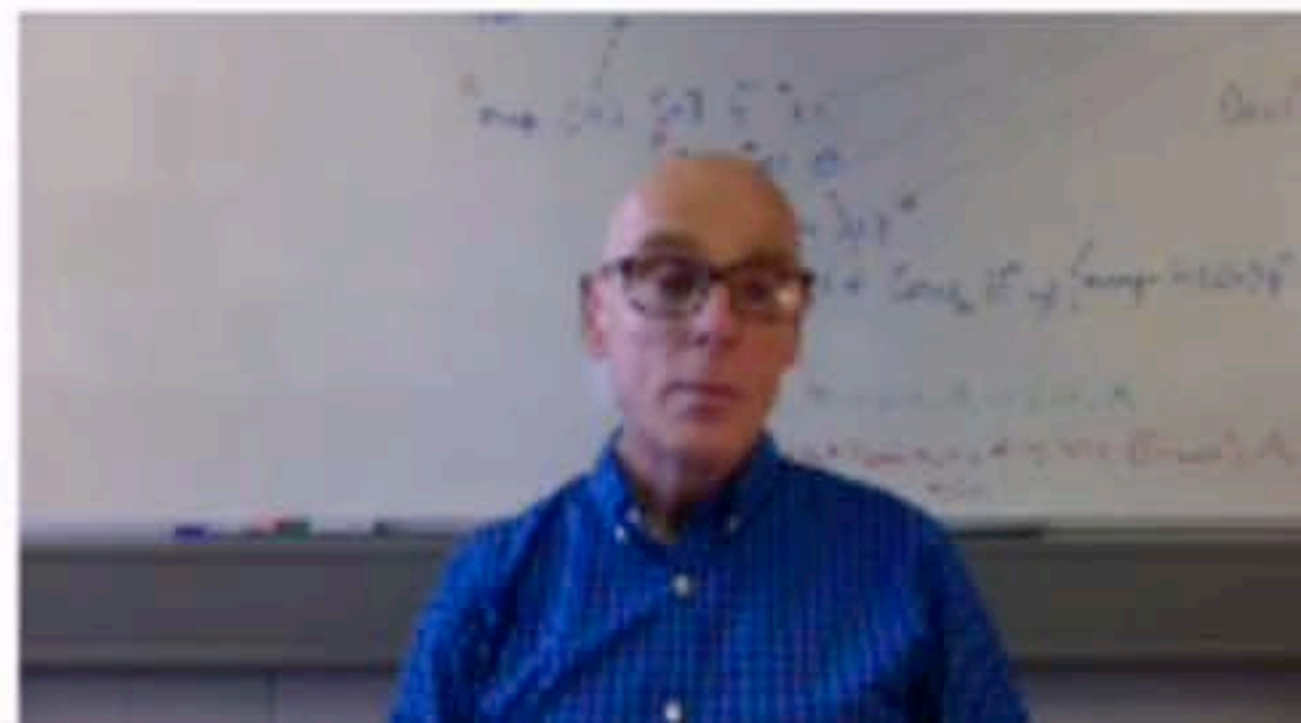
One last word ...

Just as in the case statement ... in a `receive` it's possible to pattern match against bound variables.

For example, if `Pid` is already bound then

```
receive  
  {Pid,Msg} -> ...  
end
```

will only match messages with that value as first component.



University of
Kent

University of
Kent

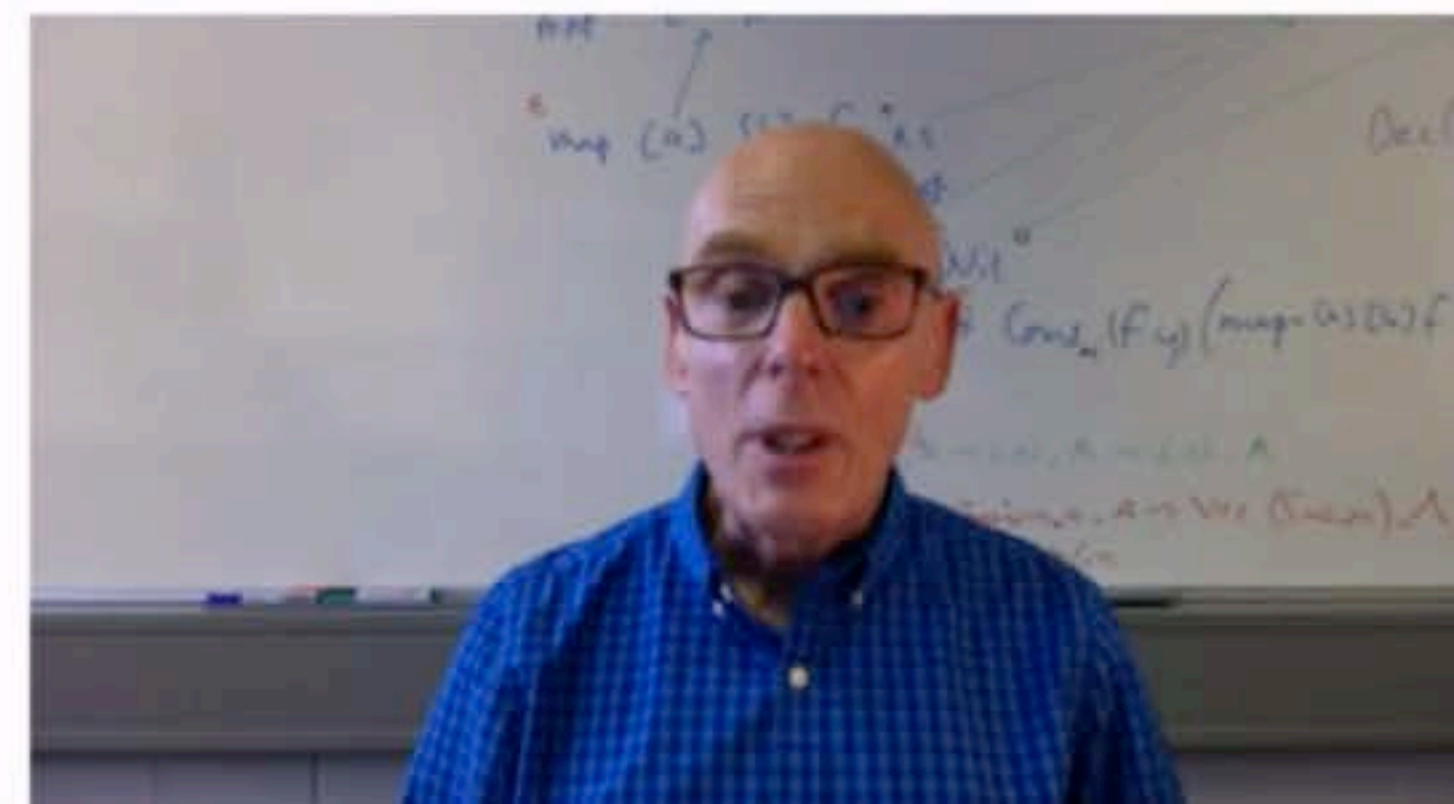
Timing and message ordering

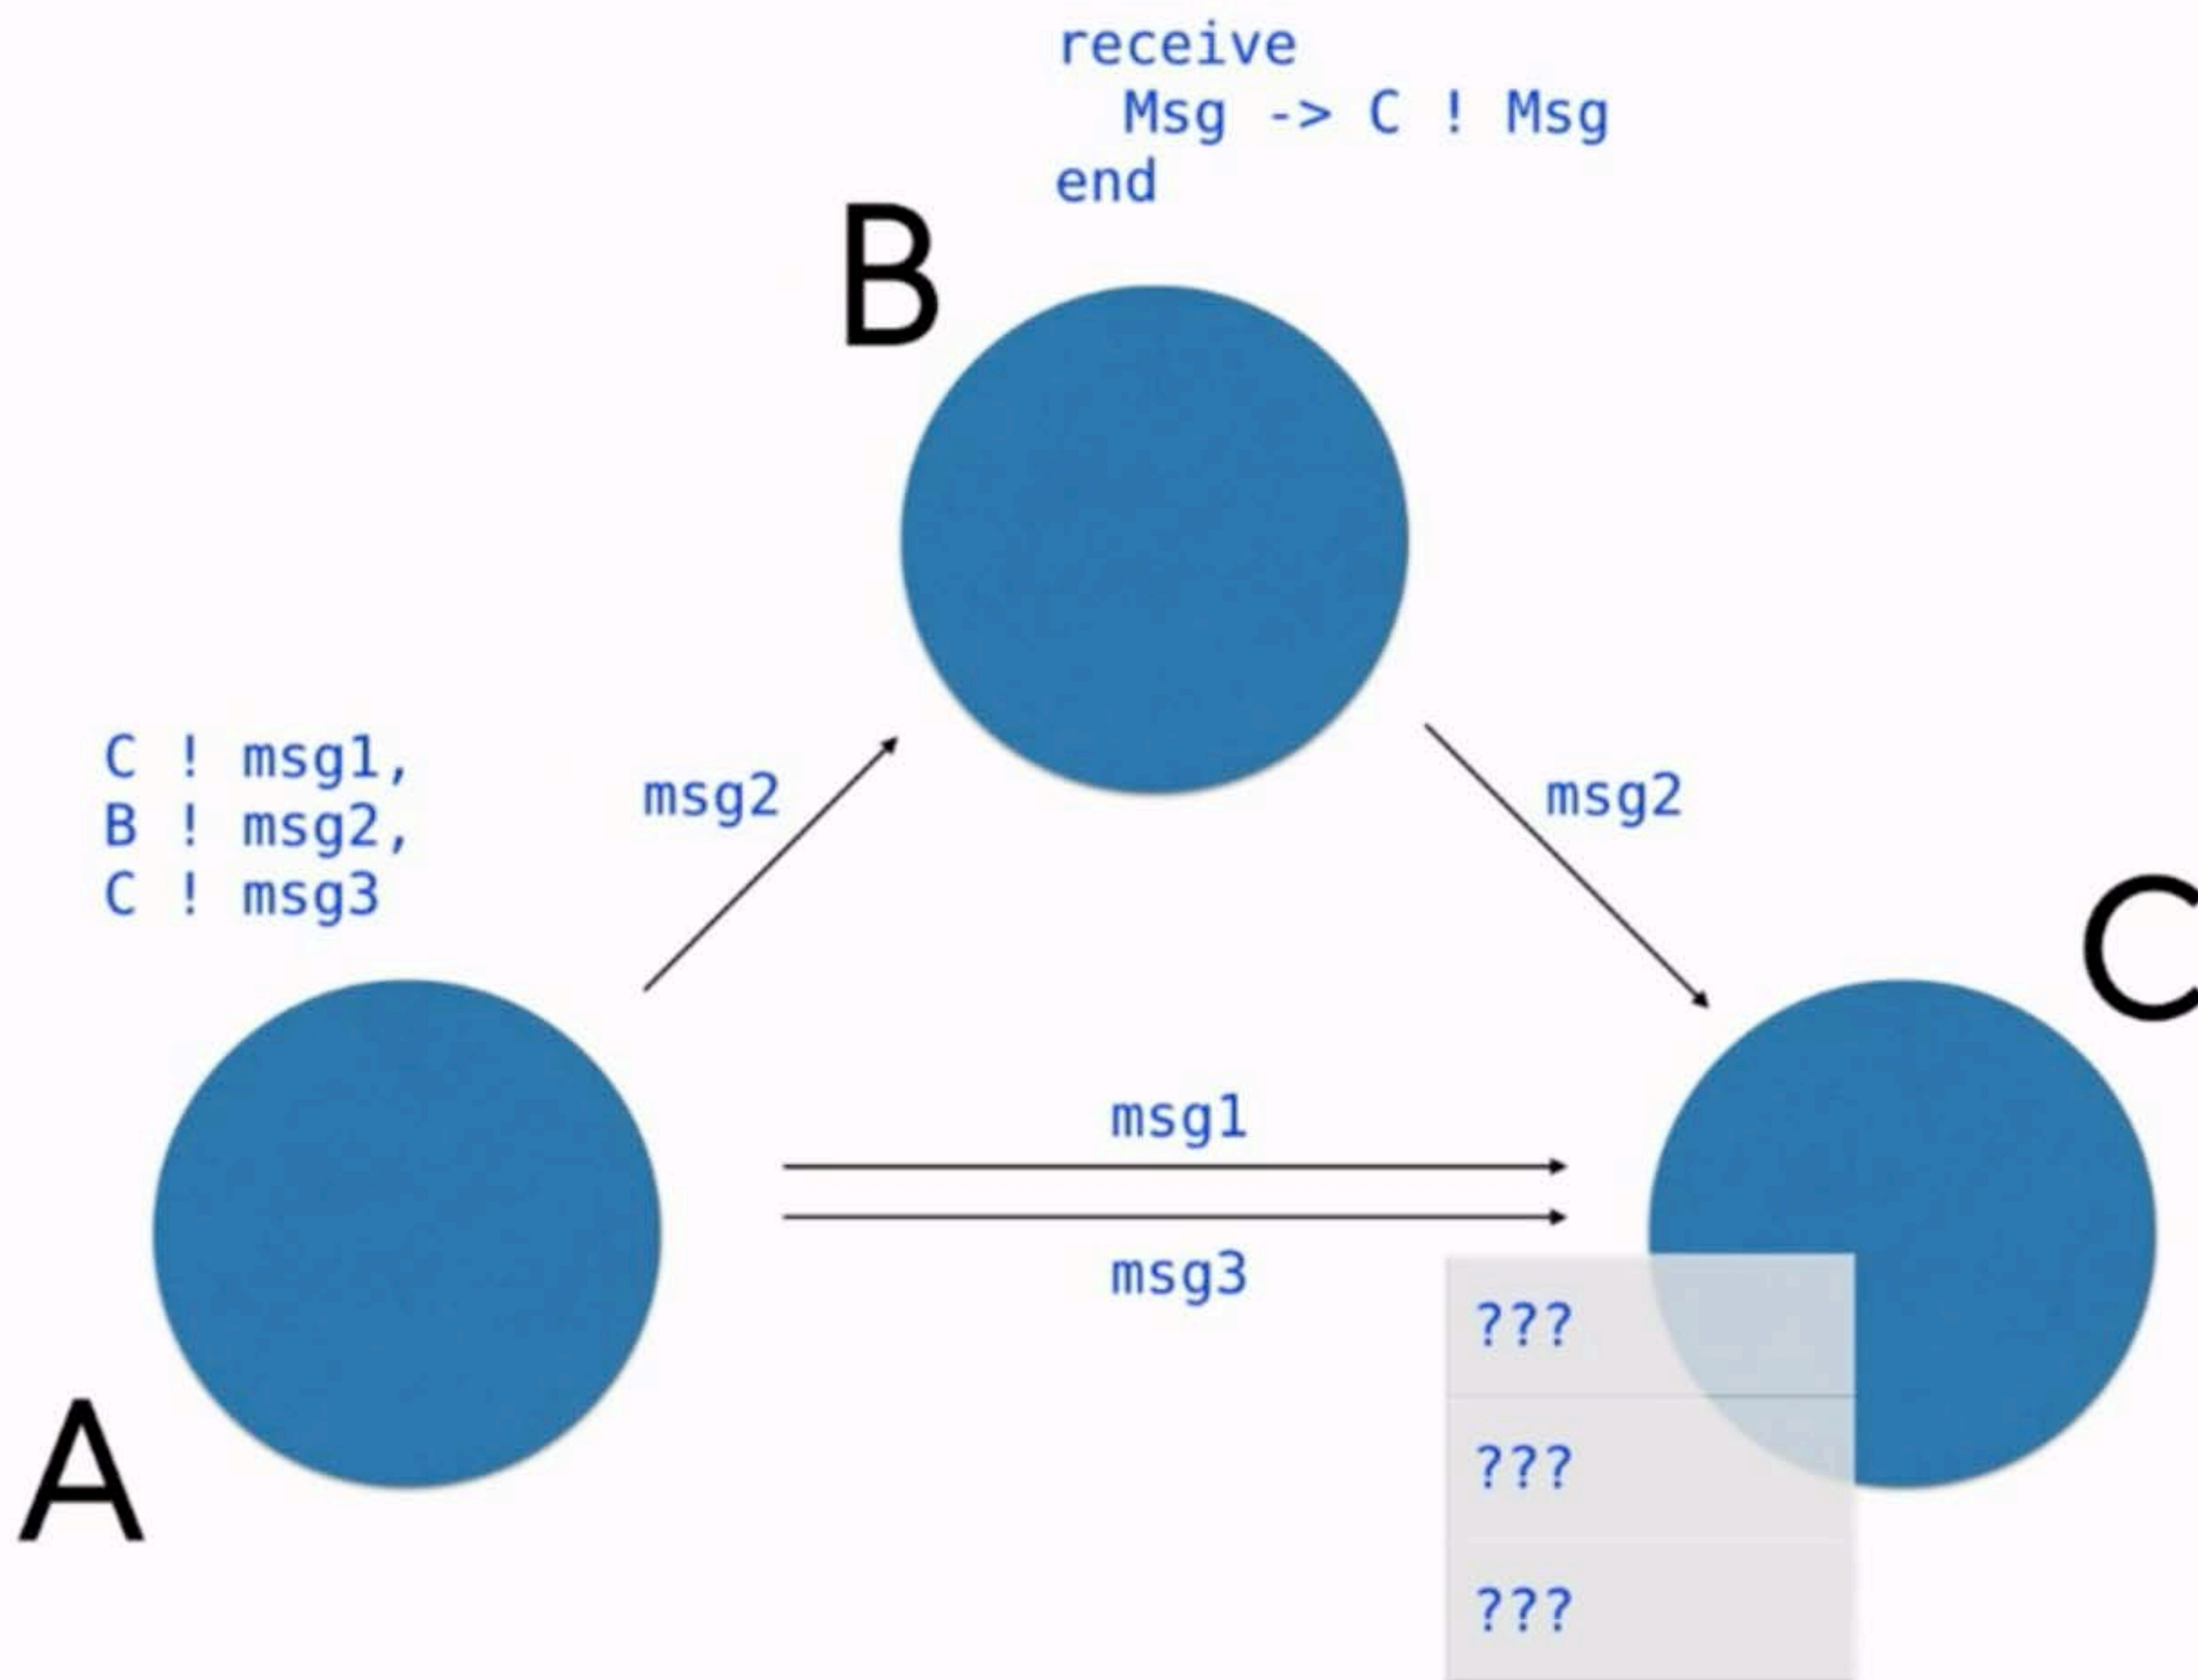


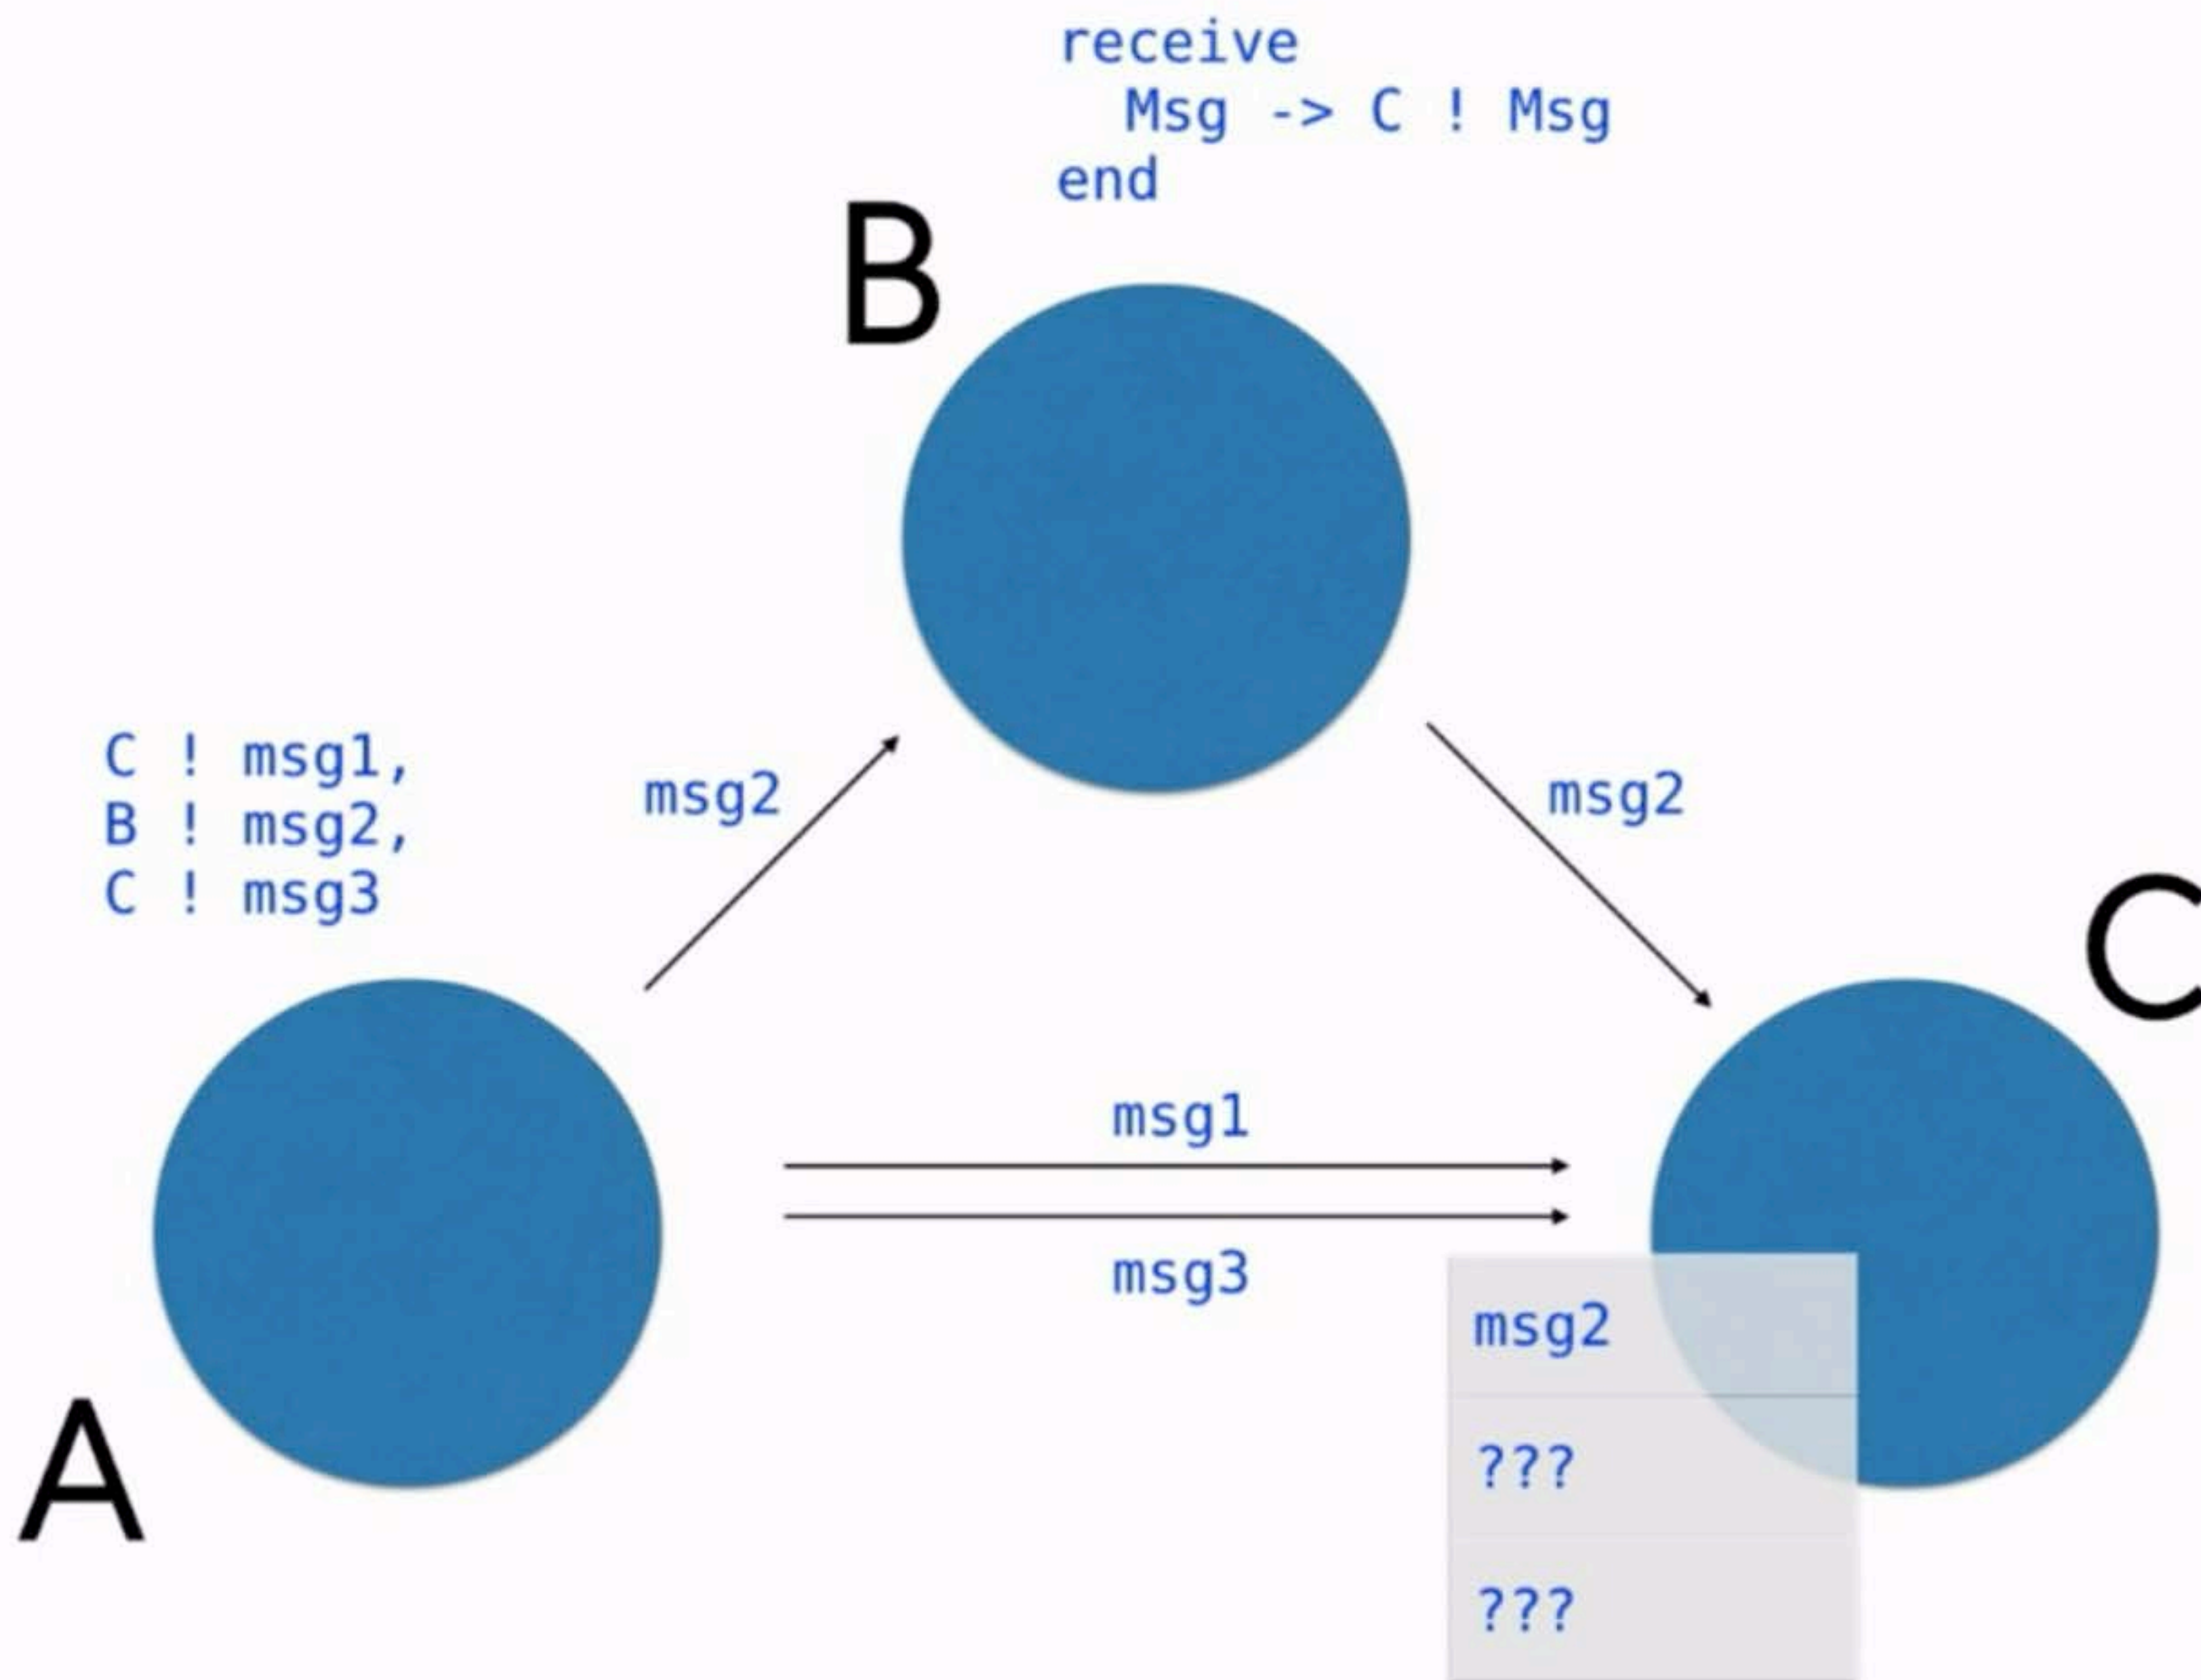
Erlang promise on message ordering

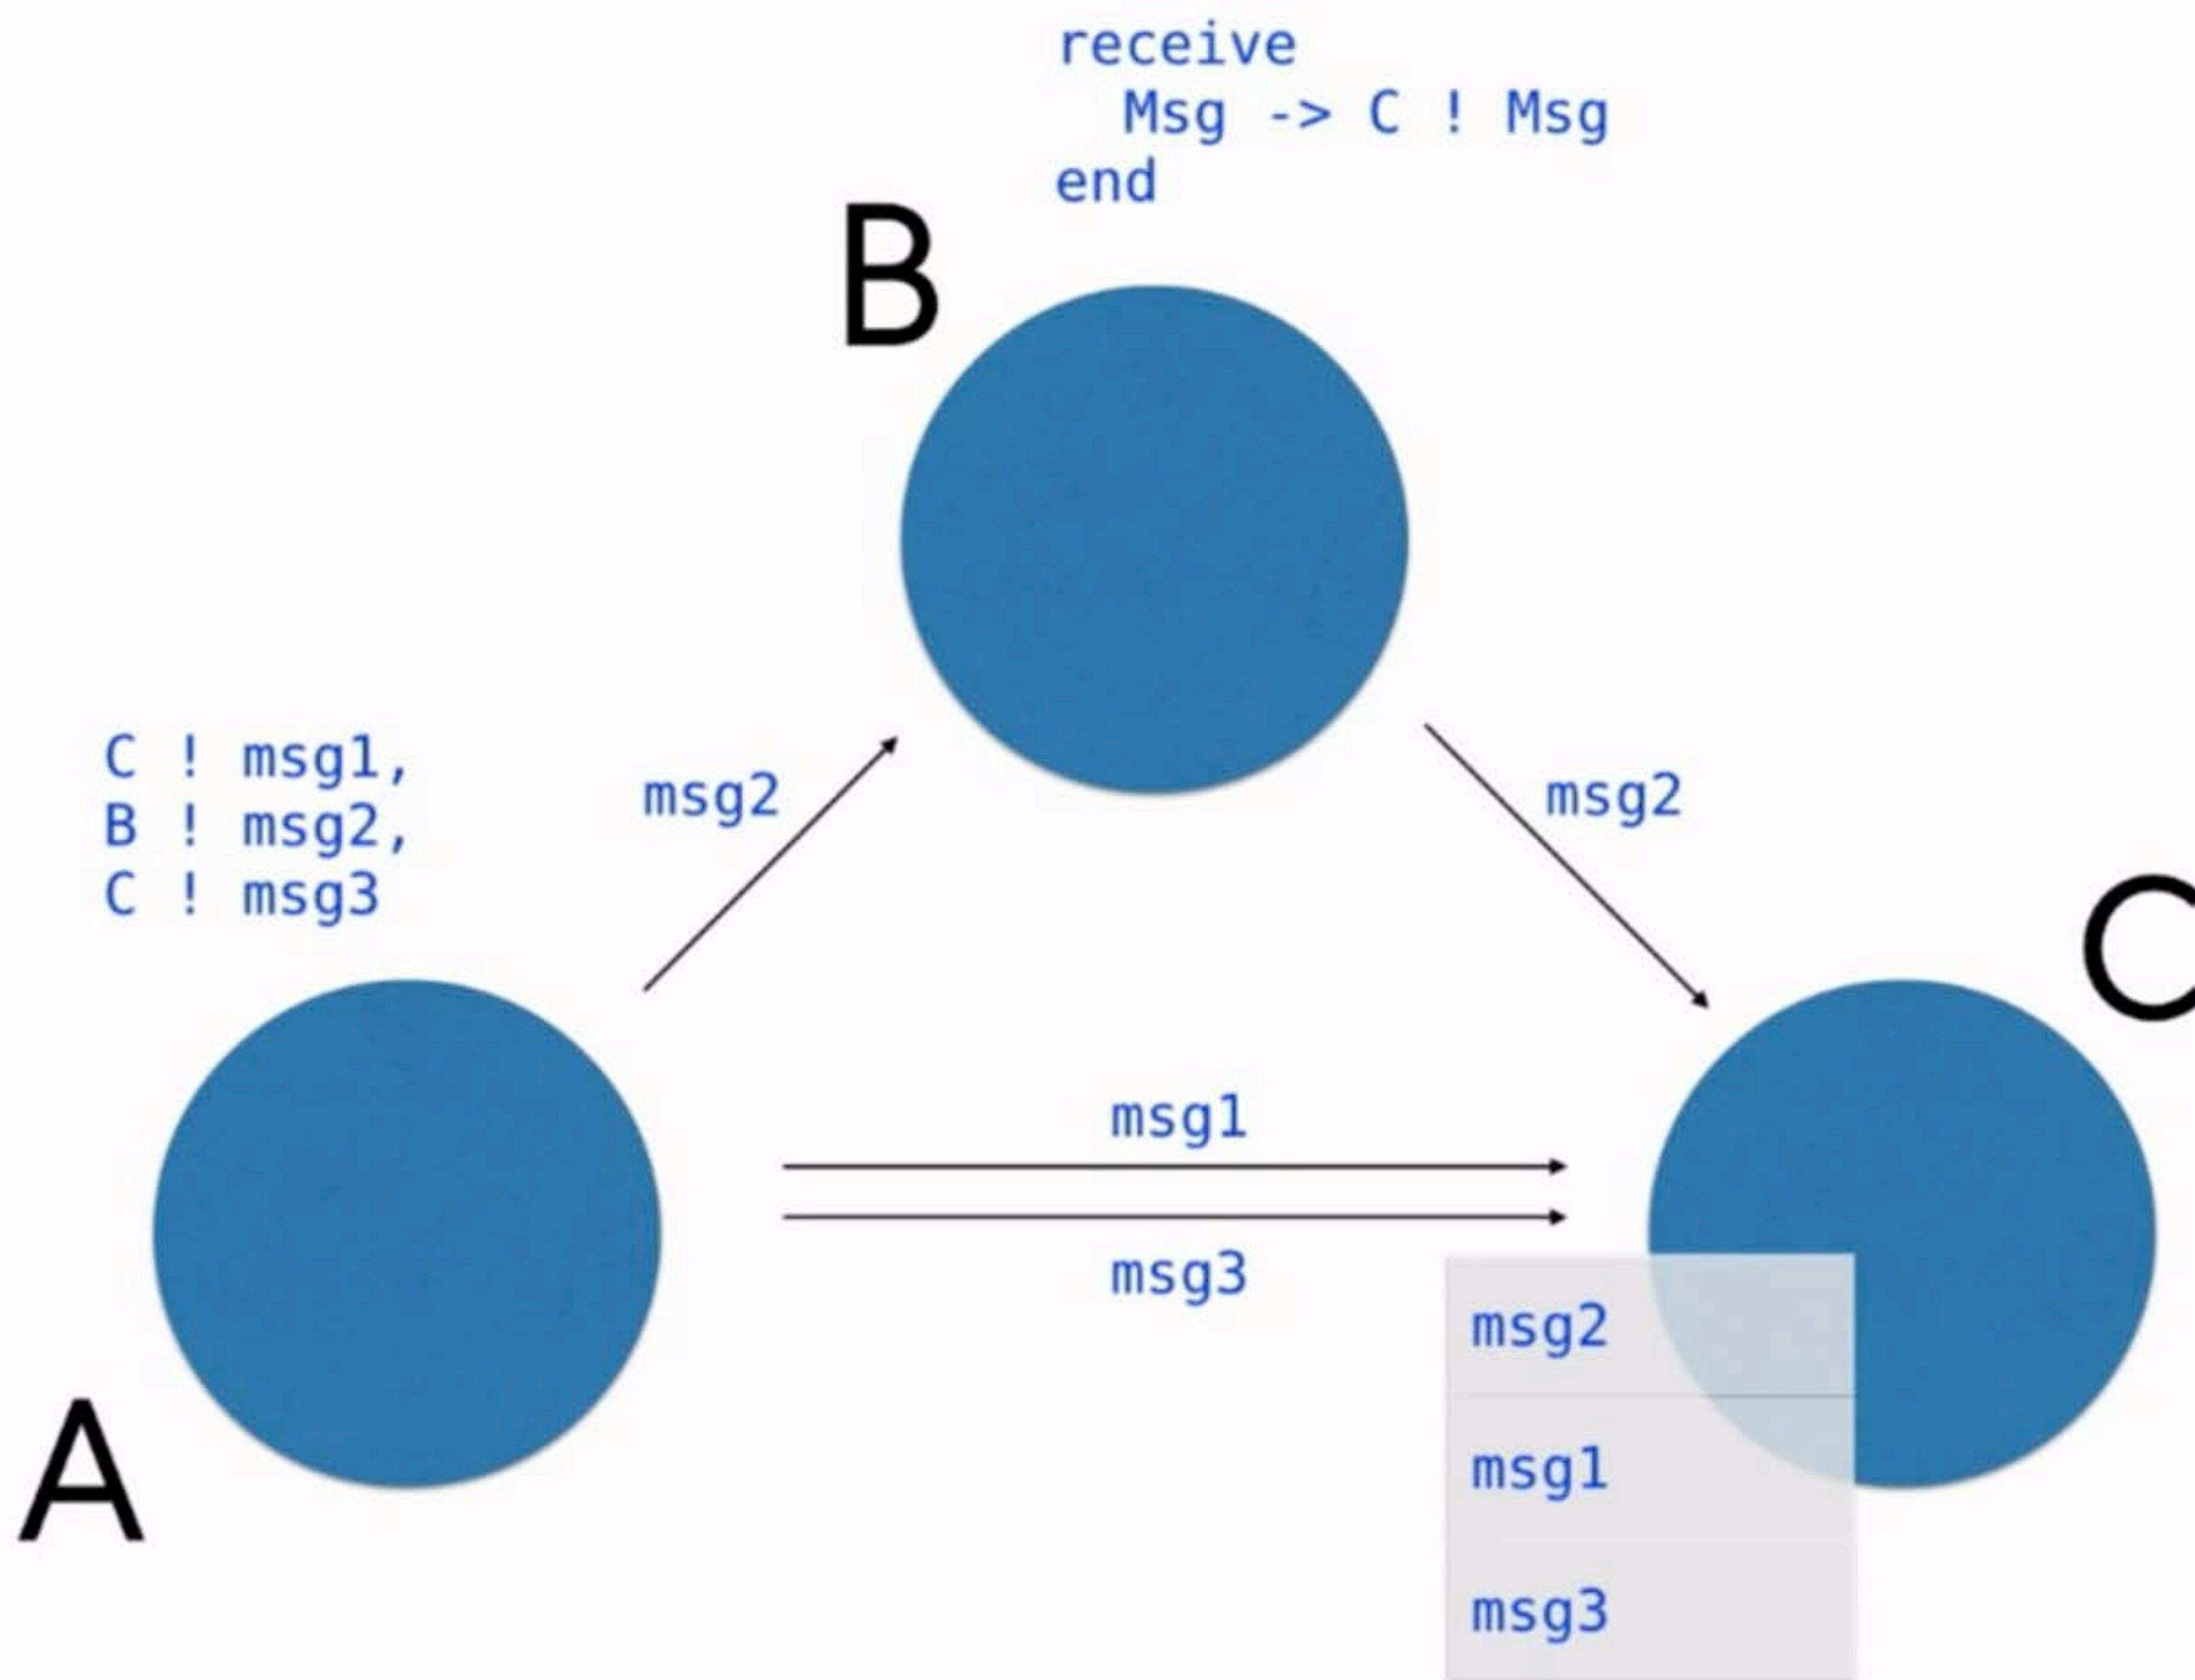
The messages sent between two particular processes will be delivered into the mailbox in the same order that they are sent.

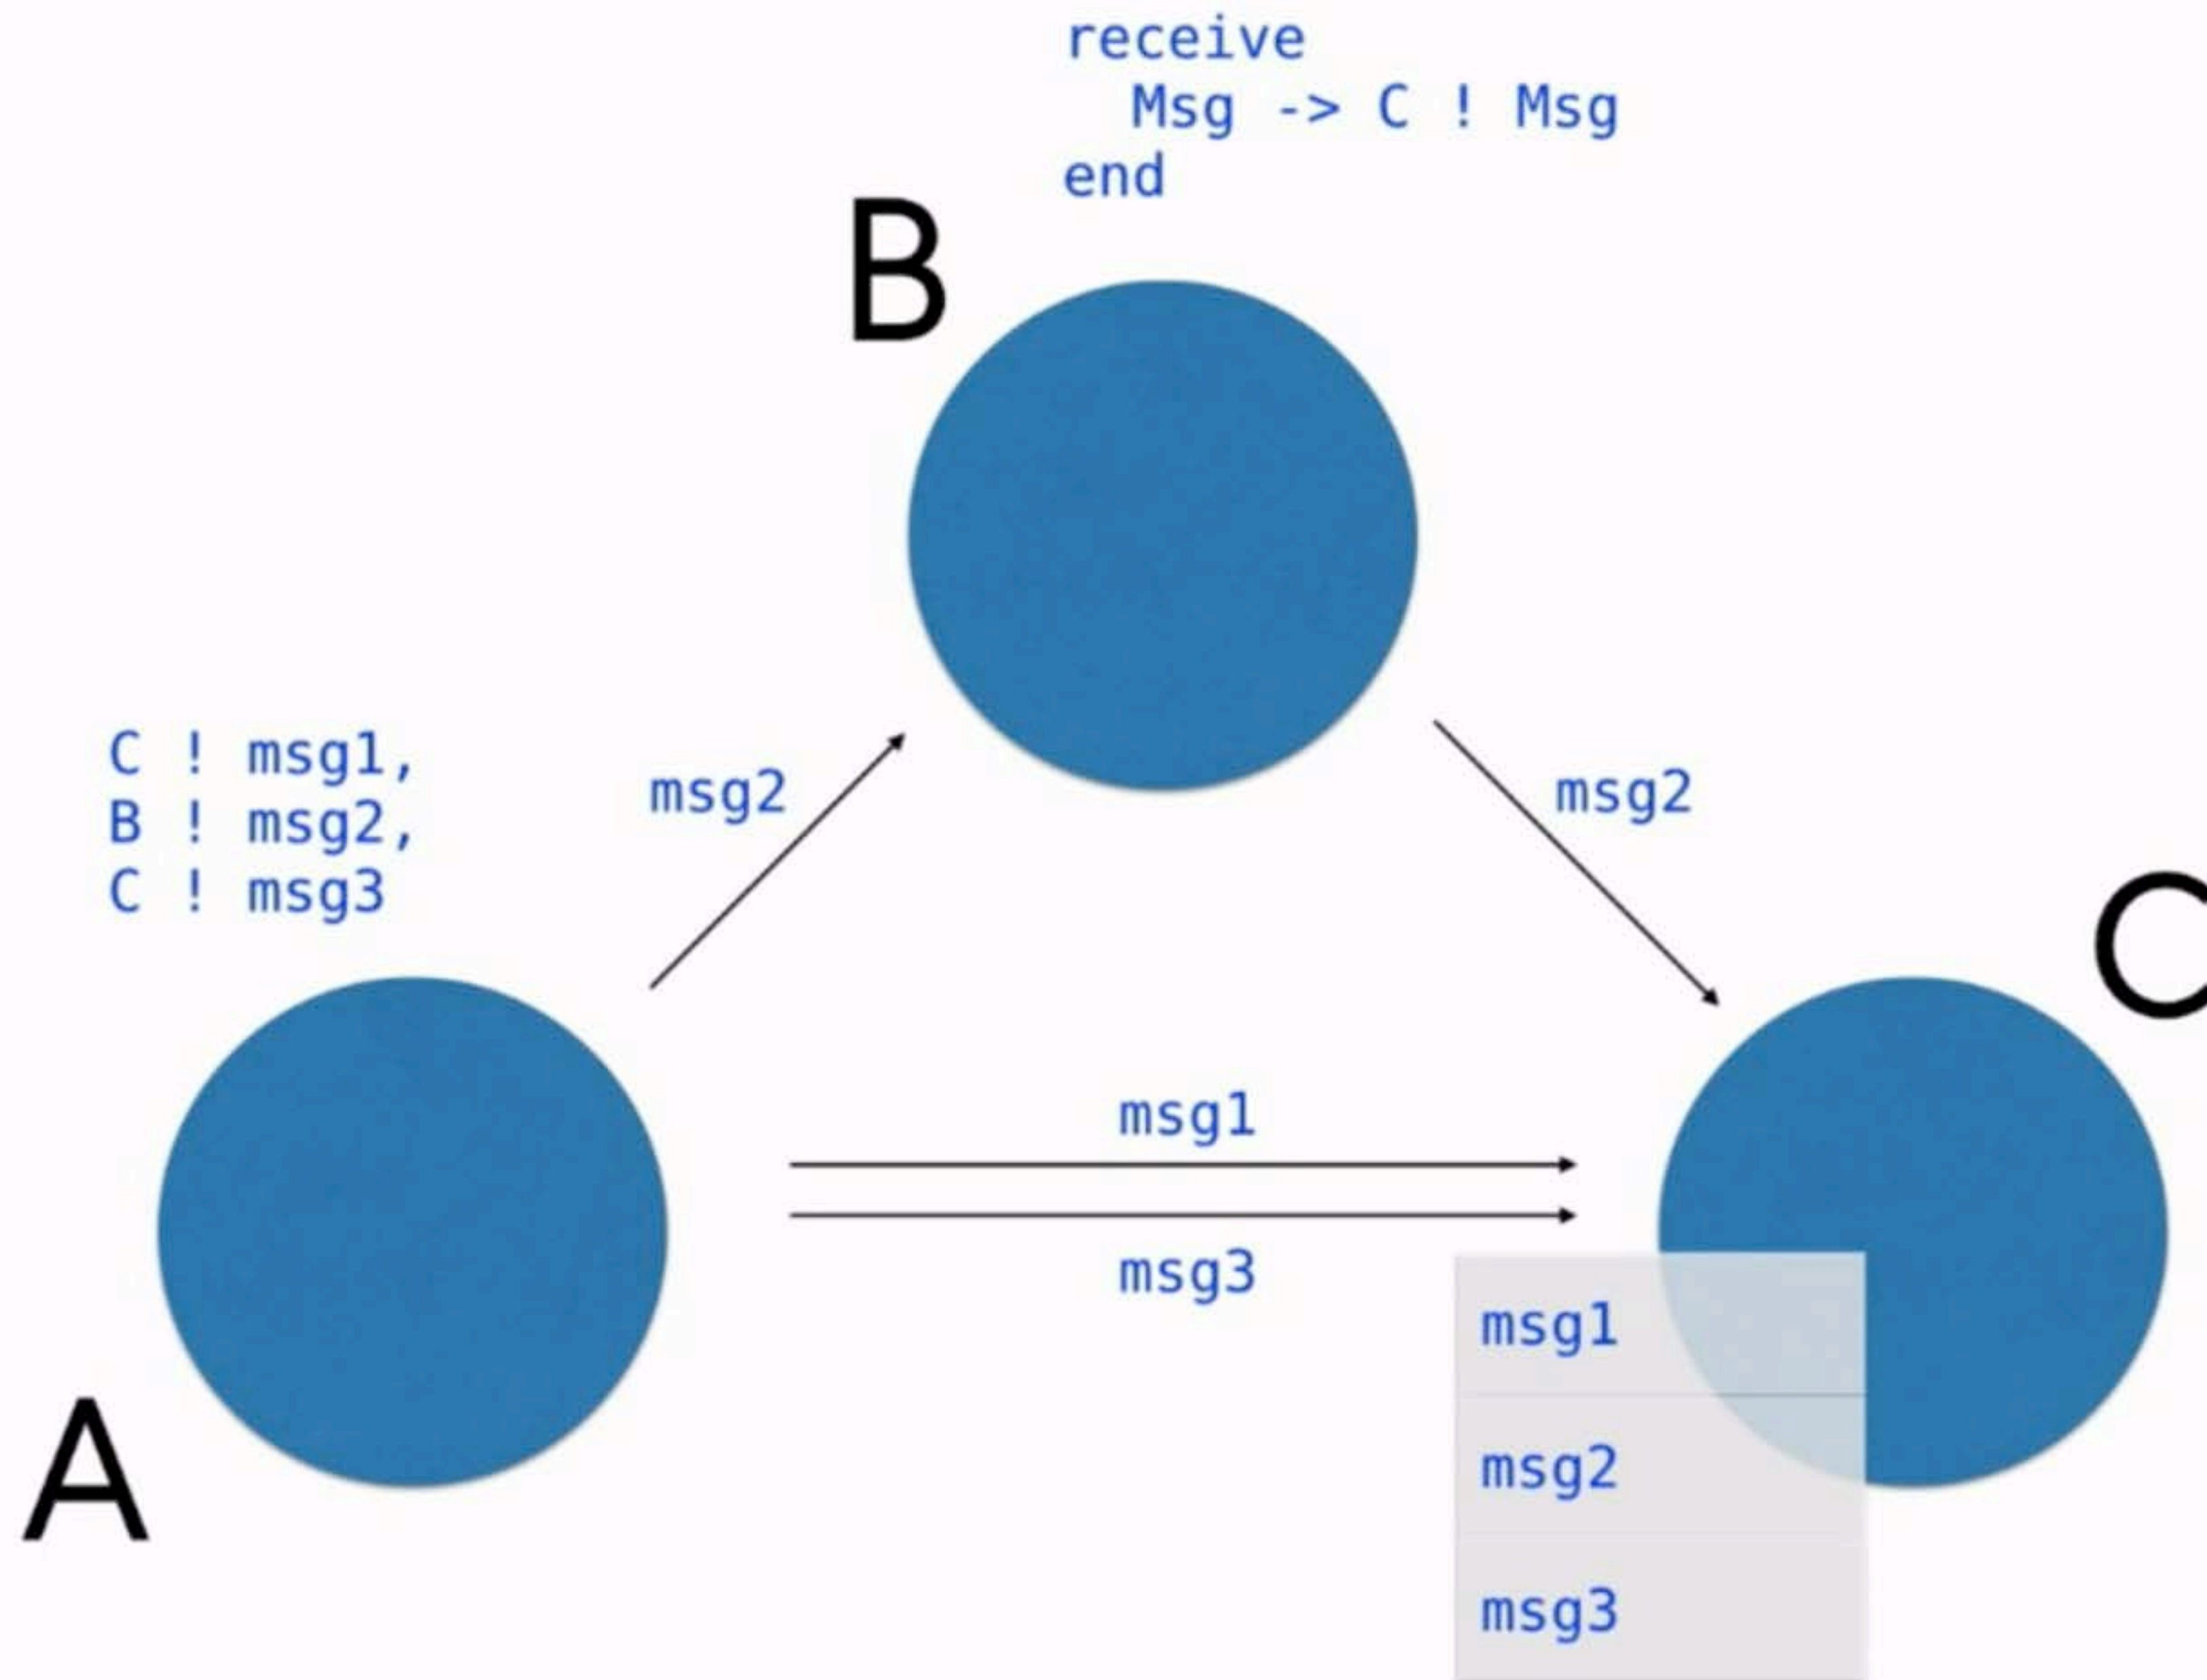
That doesn't cover indirect messaging ...

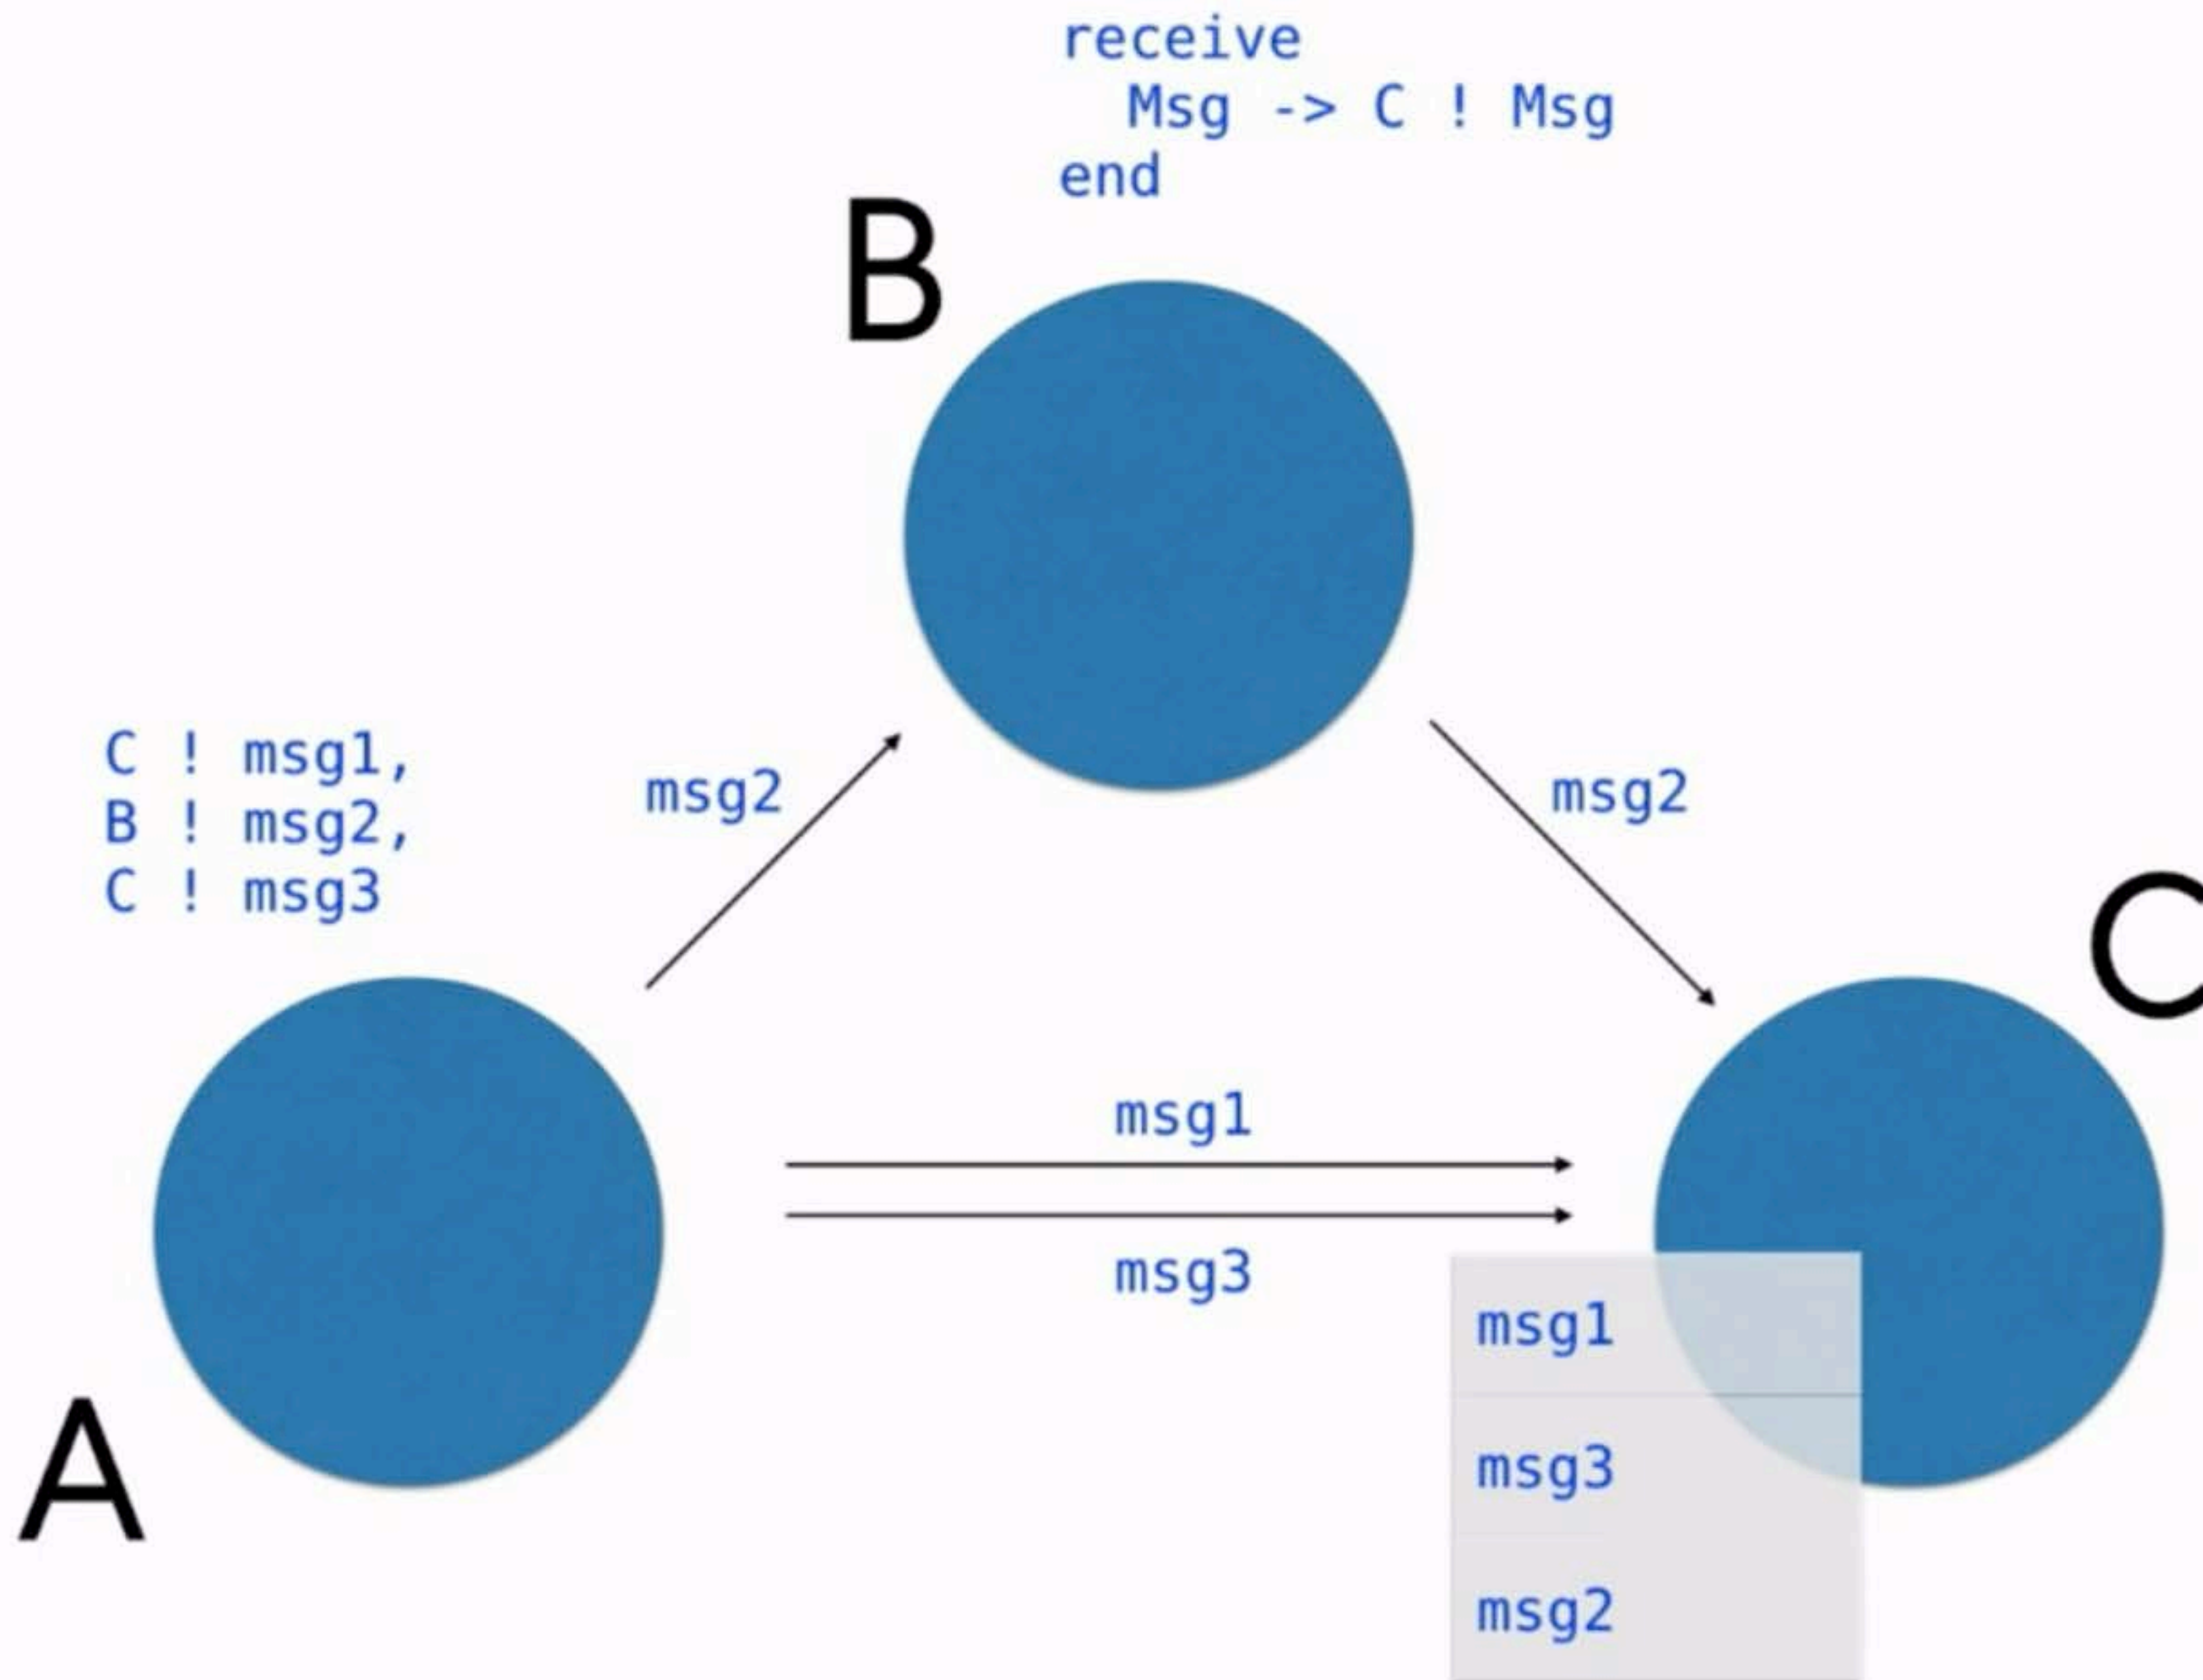










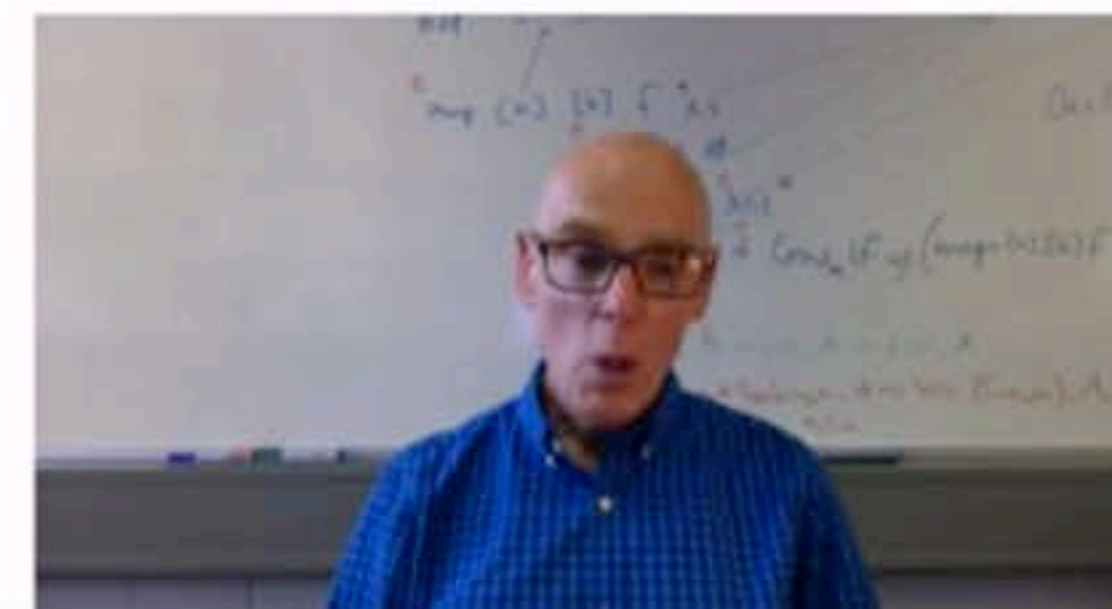


Why should this be?

Recall that on a single processor, at least, processes are scheduled in and out, allowing them to do some work – measured by function reductions.

This has the effect of sequentialising things that might be more concurrent were the resources to be available.

Real behaviour can therefore be less determinate than might be expected.



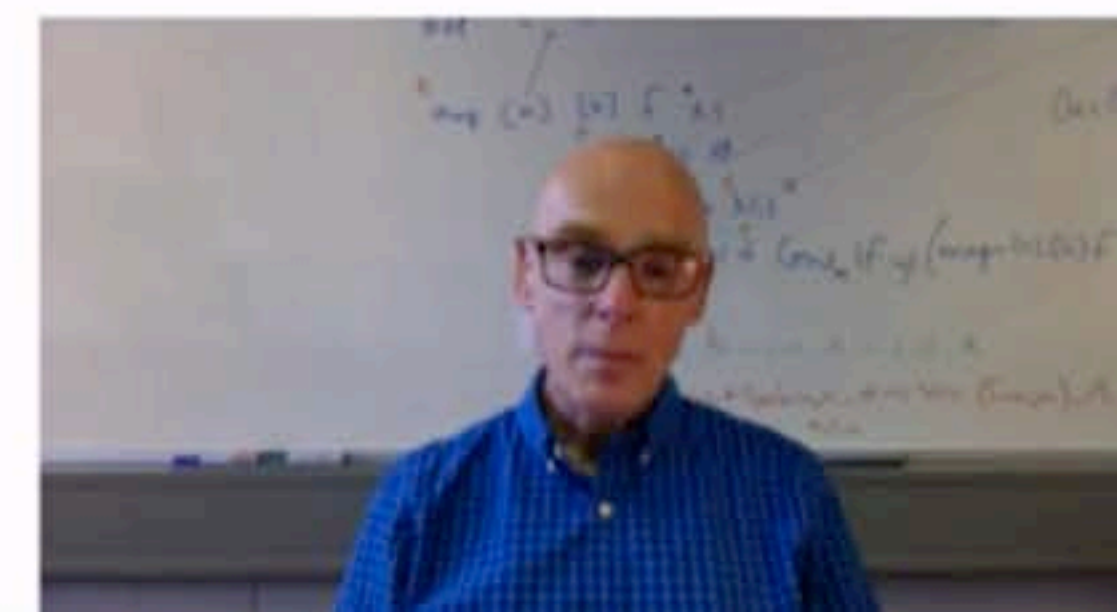
Why is this a good thing?

When we come to look at distributed Erlang, we'll see that it has exactly the same behaviour.

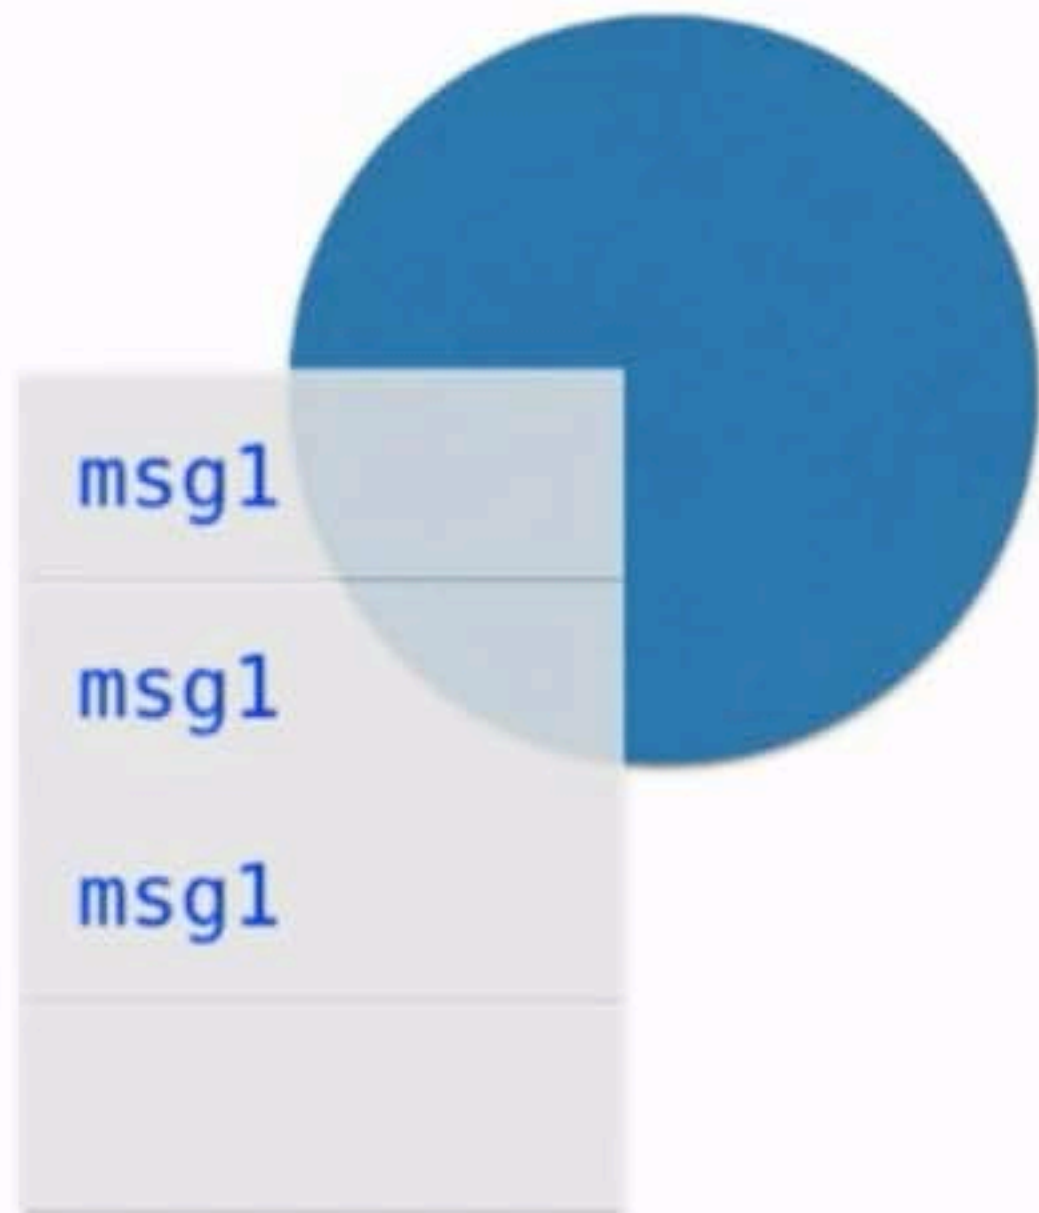
We don't distinguish between local and remote processes.



Timeouts



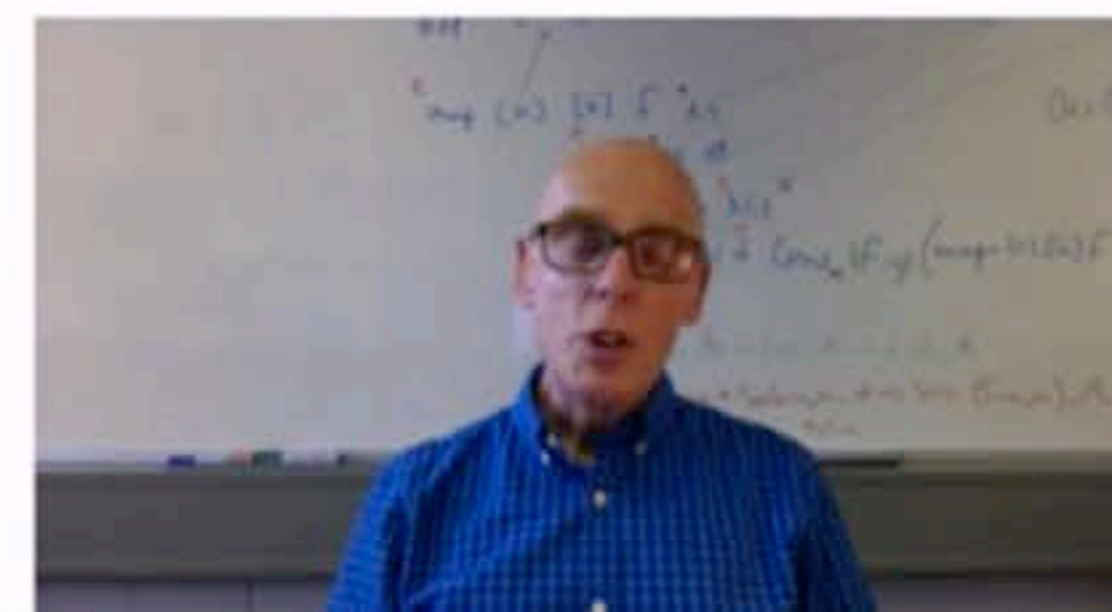
```
receive  
  msg2 -> ...;  
  msg3 -> ...;  
  msg4 -> ...  
end
```



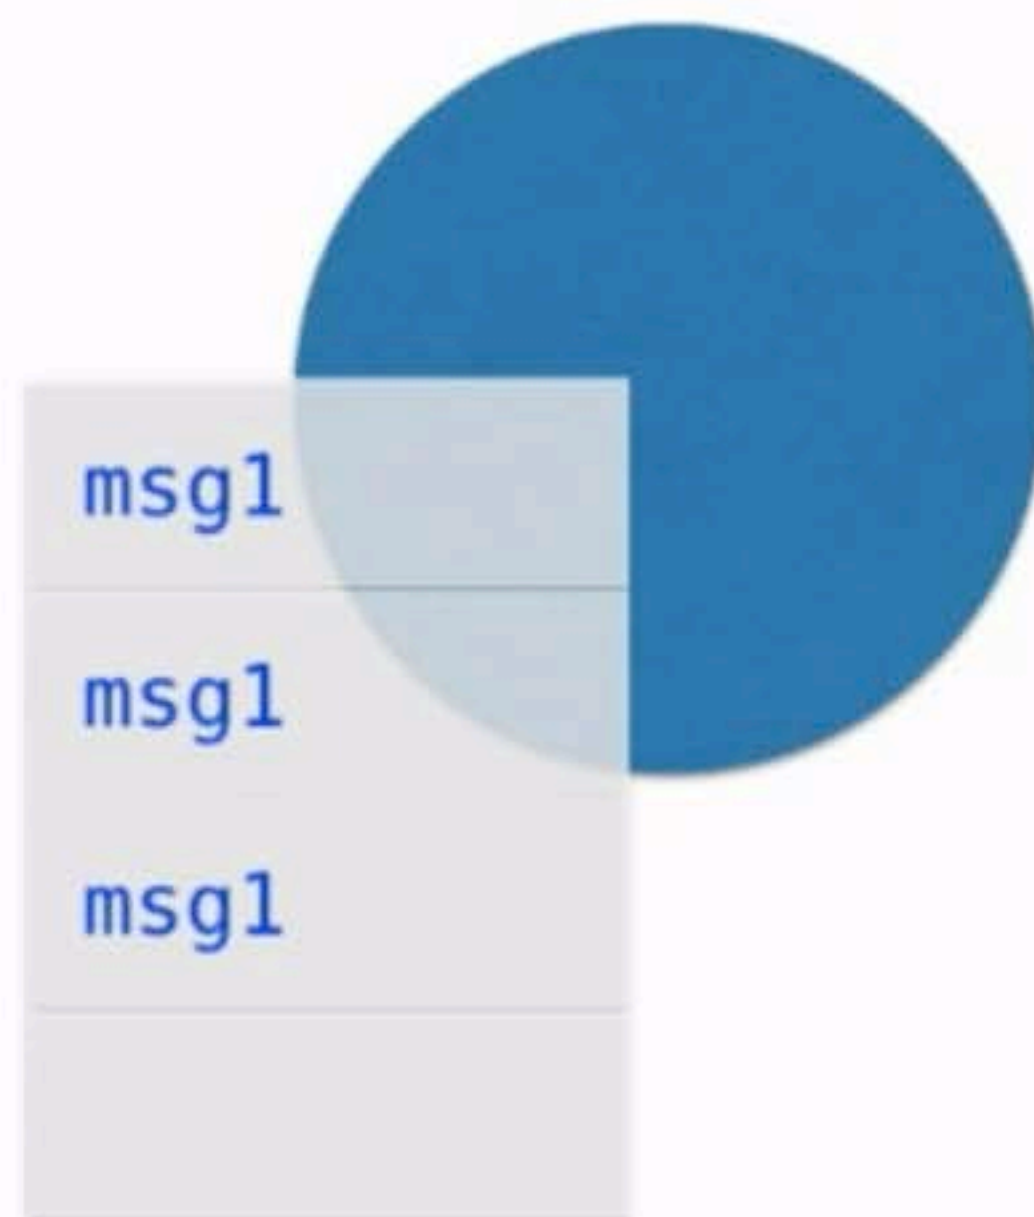
msg1

msg1

msg1



```
receive  
  msg2 -> ...;  
  msg3 -> ...;  
  msg4 -> ...  
after 500 -> ...  
end
```



Deadlock avoidance – timeout

If none of the patterns can be matched, then after waiting 500 milliseconds perform the action following the `after 500` clause.

What is the potential problem with this?

The message(s) may arrive late, and be in the mailbox later in the program.



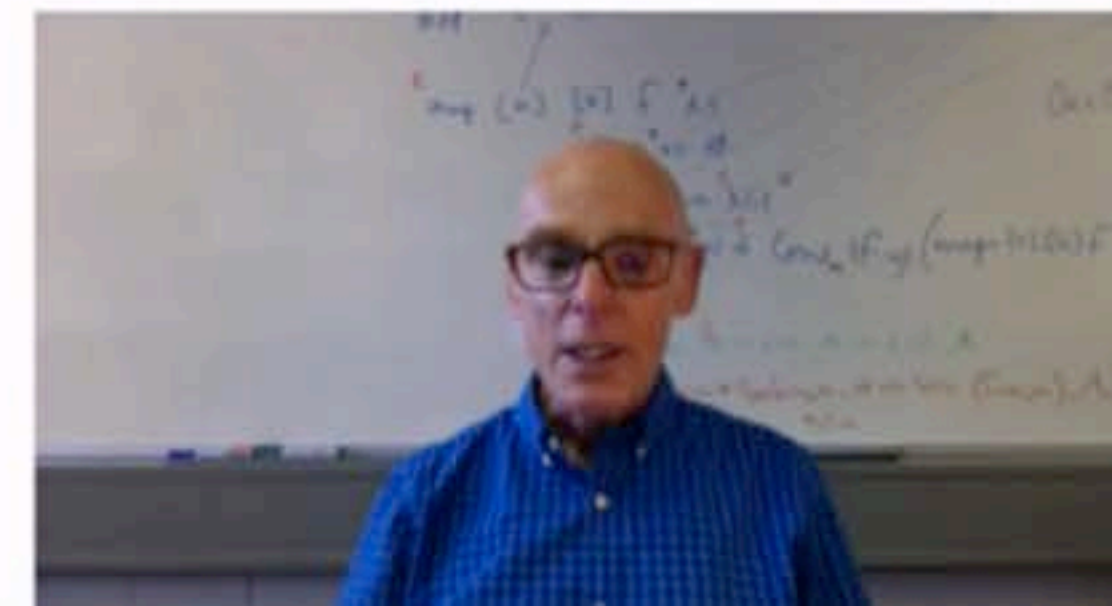
Flushing the mailbox

How can we remove all the messages that are in the mailbox?

```
clear() ->  
  receive ->  
    Msg -> ok  
  end,  
  clear().
```

What is the problem with this?

It never terminates ... it waits for messages, reads and repeats!



Flushing the mailbox

How can we remove all the messages that are in the mailbox?

```
clear() ->  
  receive  
  _Msg -> clear()  
after 0 ->  
  ok  
end.
```

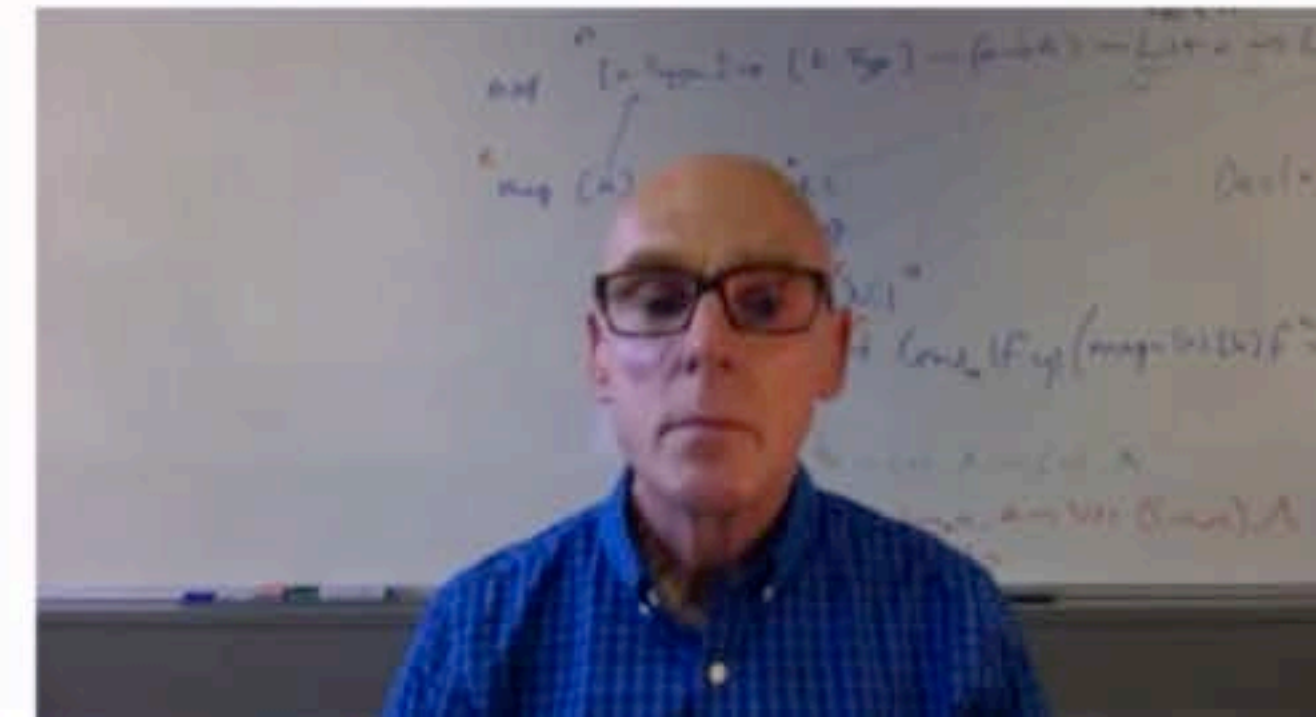
If a message is present, read it and repeat ... if not (`after 0`) then terminate.



The logo of the University of Kent, featuring the text "University of Kent" in a blue serif font. The word "University of" is in a smaller size and positioned above the word "Kent". The entire logo is centered within a light gray rectangular box.

University of
Kent

Establishing communication: Pids and named processes

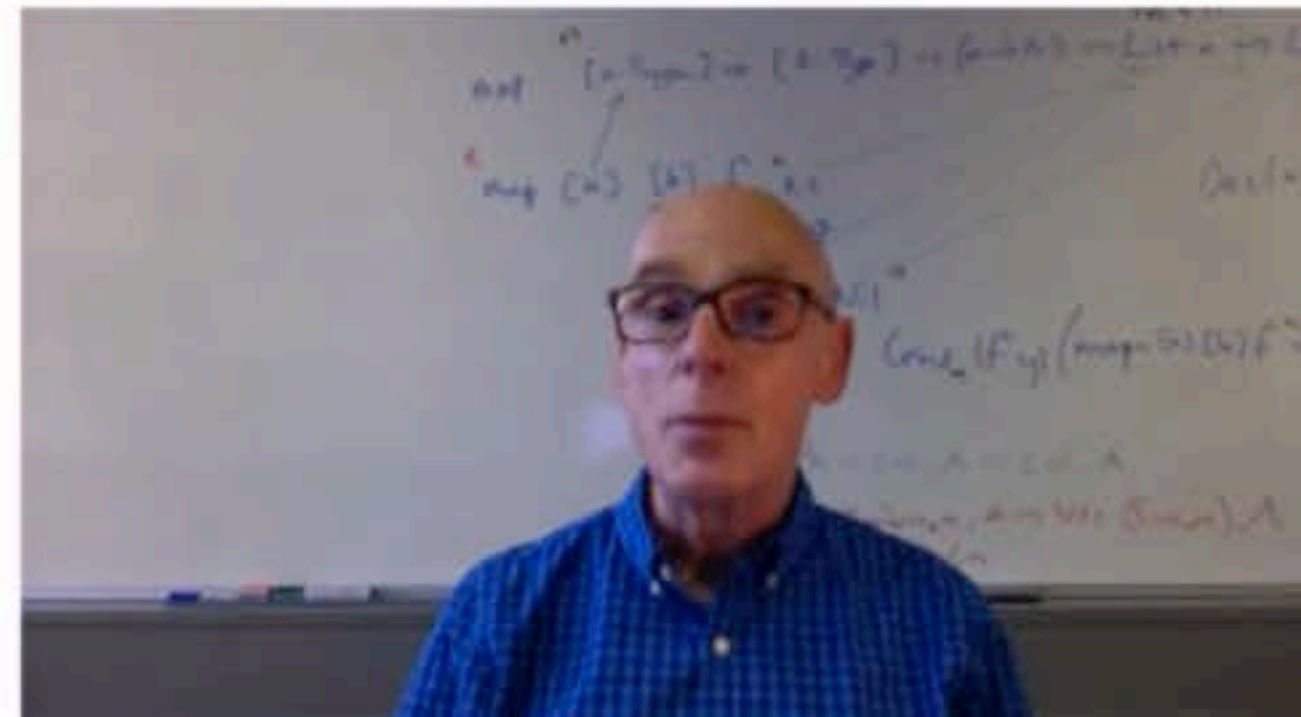


Sending messages

`Pid ! Message` will send the `Message` to the process with id `Pid`.

For us to be able to do this, we need to know the `Pid` ...

... and so this needs to be communicated to us somehow.



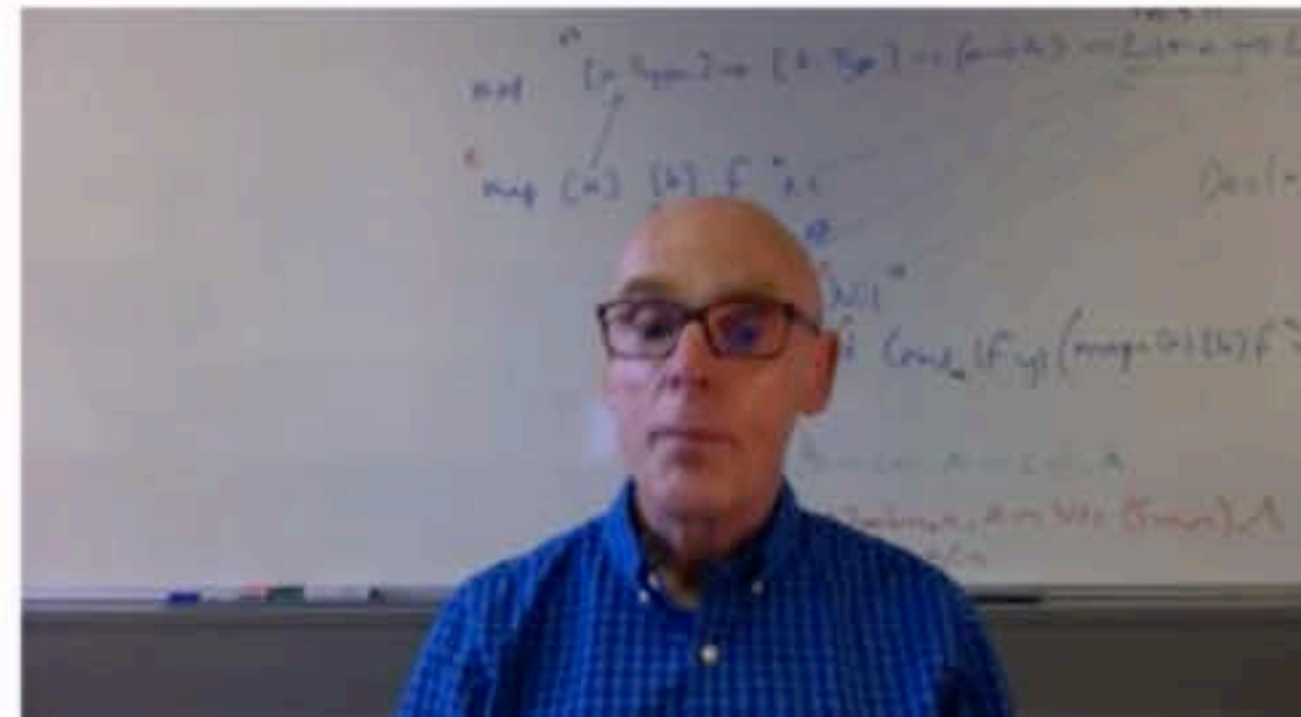
Include your `Pid`

The typical pattern is to send a message to a (`Server`) process with your `Pid` ...

...and your recipient can then send a message back to you.

```
Server ! {self(), ping},  
receive  
  pong -> ...  
end
```

```
receive  
  {Pid, ping} ->  
    Pid ! pong  
end
```



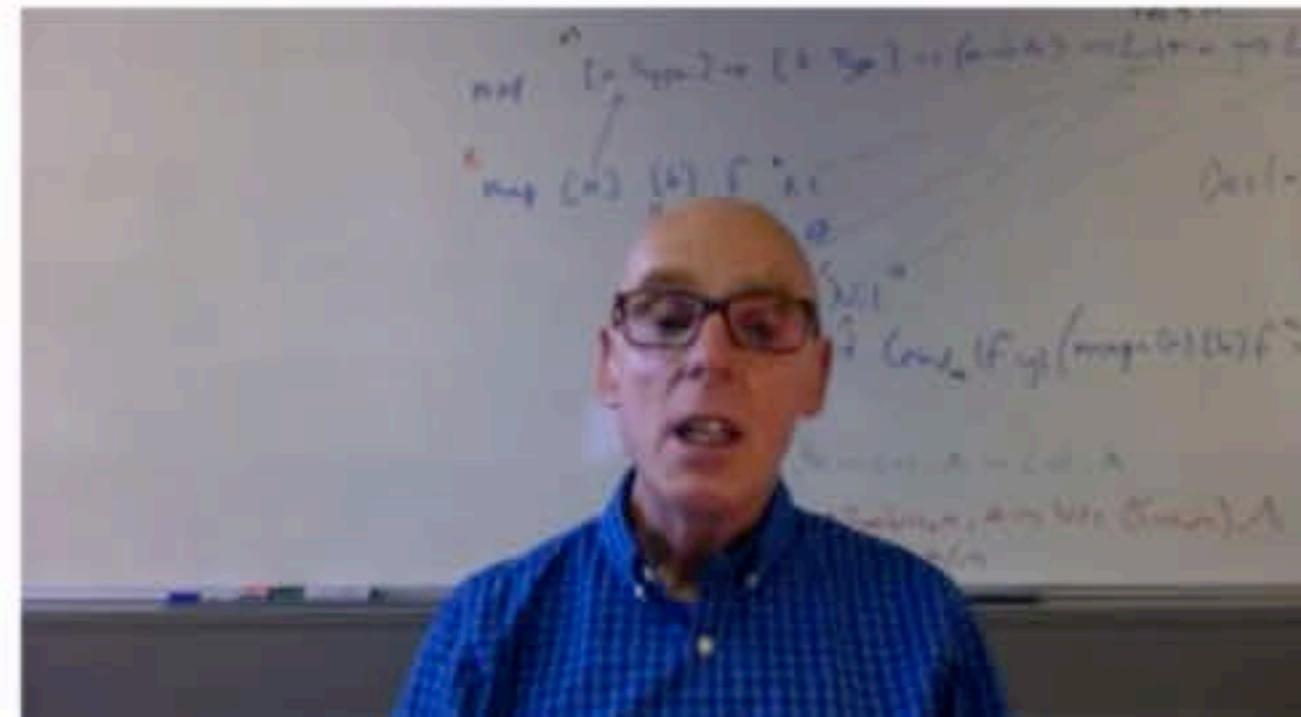
Include your `Pid`

It's also possible to include a `Pid` in the return message ...

... and the client can check that the message has come from the right place.

```
Server ! {self(), ping},  
receive  
  {Server, pong} -> ...  
end
```

```
receive  
  {Pid, ping} ->  
    Pid ! {self(), pong}  
end
```



Named processes

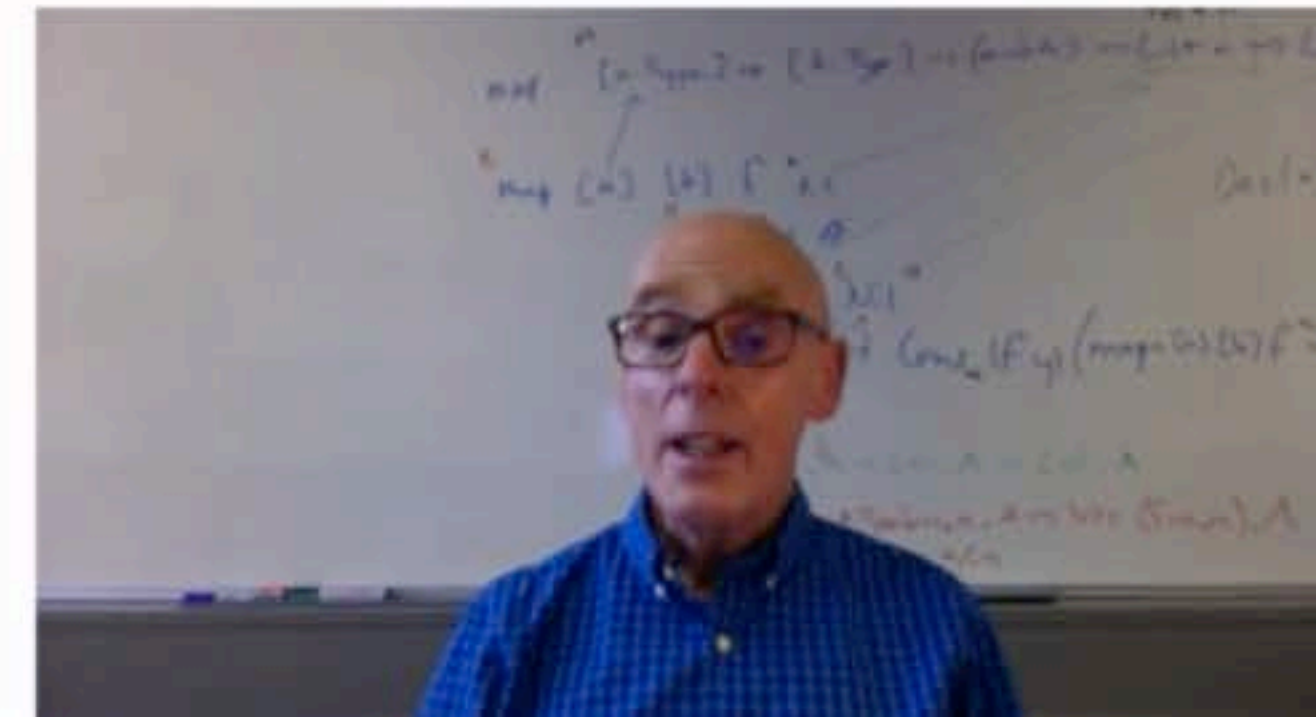
As an alternative we can name a process, e.g. `server` ...

... and send messages to that name: `server ! Msg`.

Typically `spawn` and `register` are done in one nested call ...

```
Server = spawn(M,F,[]),  
register(server,Server),  
server ! {self(), ping}
```

```
register(server,spawn(M,F,[])),
```



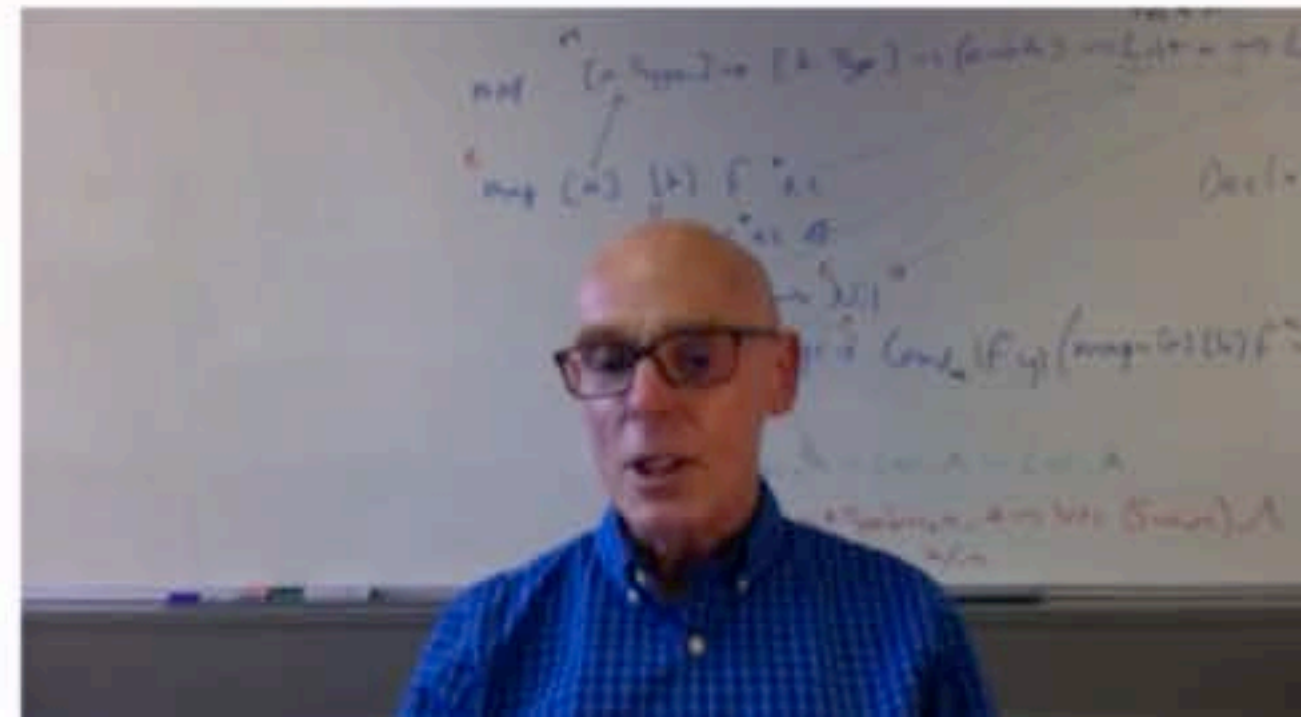
Names *versus* Pids

Sending a message to a non-existent name is an *error* ...

... the assumption is that named processes should not disappear.

So, typically we name "static", long-lived processes.

A convention is that the named process gets the name of the module where it is defined ... this assumes only one per module.



A hint of things to come

Suppose that we spawn the server process, and name it `server` ...

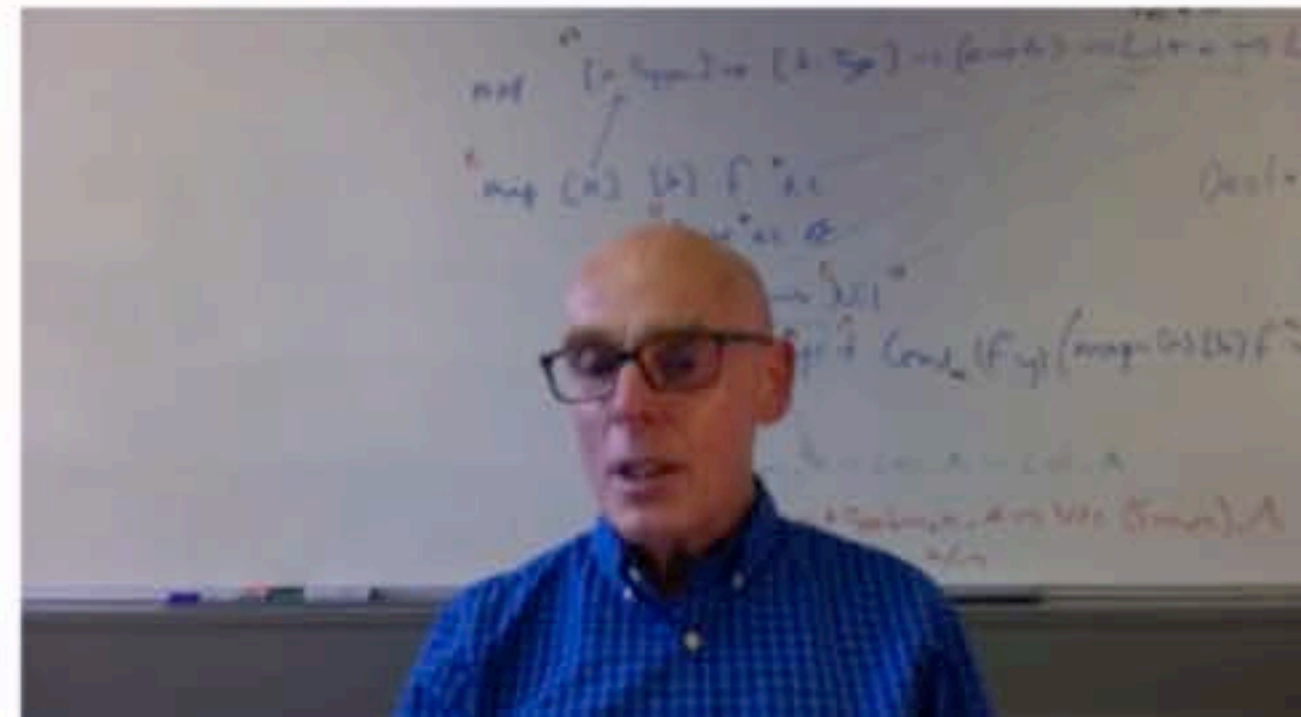
... and other processes can then send messages to it: `server ! ping`.

What if the `ping` is sent before registration is complete?

This is a *race condition*.

```
Server = spawn(M,F,[]),  
register(server,Server),
```

```
server ! ping
```



University of
Kent